Predictive Systems: Living with Imperfect Predictors

ĽUBOŠ PÁSTOR and ROBERT F. STAMBAUGH*

ABSTRACT

We develop a framework for estimating expected returns—a *predictive system*—that allows predictors to be imperfectly correlated with the conditional expected return. When predictors are imperfect, the estimated expected return depends on past returns in a manner that hinges on the correlation between unexpected returns and innovations in expected returns. We find empirically that prior beliefs about this correlation, which is most likely negative, substantially affect estimates of expected returns as well as various inferences about predictability, including assessments of a predictor's usefulness. Compared to standard predictive regressions, predictive systems deliver different expected returns with higher estimated precision.

MANY STUDIES IN FINANCE analyze comovement between expected asset returns and various observable quantities, or "predictors." A question of frequent interest is how x_t, a vector of predictors observed at time t, is related to μ_t, the conditional expected return defined in the equation

$$r_{t+1} = \mu_t + u_{t+1}, \tag{1}$$

where r_{t+1} denotes the stock return from time t to time $t + 1$, and the unexpected return u_{t+1} has mean zero conditional on all information available at time t. One approach to modeling expected returns is to use a "predictive regression" in which r_{t+1} is regressed on x_t and the expected return is given by $\mu_t = a + b'x_t$, where a and b denote the regression's intercept and slope coefficients.[1] This

[1] Of the many studies that estimate predictive regressions for stock returns, some early examples include Fama and Schwert (1977), Rozeff (1984), Keim and Stambaugh (1986), Campbell (1987), and Fama and French (1988). There is also a substantial literature analyzing econometric issues

approach seems too restrictive in modeling expected return as an exact linear function of the observed predictors. It seems more likely that the predictors are imperfect, in that they are correlated with μ_t but cannot deliver it perfectly.

At the same time, the predictive regression approach seems too lax in ignoring what we argue is a likely property of the unexpected return—its negative correlation with the innovation in the expected return. For example, if the expected return obeys the first-order autoregressive process

$$\mu_{t+1} = (1 - \beta)E_r + \beta\mu_t + w_{t+1}, \tag{2}$$

then it seems likely that the correlation between the unexpected return and the innovation in the expected return is negative, or that $\rho_{uw} \equiv \rho(u_{t+1}, w_{t+1}) < 0$. That is, unanticipated increases in expected future returns (or discount rates) should be accompanied by unexpected negative returns. The likely negativity of ρ_{uw}, which is not exploited in estimating the predictive regression, emerges as an important consideration in estimating expected returns when predictors are imperfect.

We develop an approach to estimating expected returns that generalizes the standard predictive regression approach. The framework we propose, which we term a *predictive system*, allows the predictors in x_t to be imperfect, in that $\mu_t \neq a + b'x_t$. When predictors are imperfect, their predictive ability is supplemented by information in lagged returns as well as lags of the predictors, and the predictive system delivers that information via a parsimonious model. The predictive system also allows us to explore roles for a variety of prior beliefs about the behavior of expected returns, chief among which is the belief that unexpected returns are negatively correlated with innovations in expected returns ($\rho_{uw} < 0$). The correlation ρ_{uw} plays a key role in determining how the information in lagged returns and predictors is used as well as how important that information is in explaining variation in expected returns.

The additional information in lagged returns is used in an interesting way. Suppose that recent returns have been unusually low. On the one hand, one might think that the expected return has declined, since a low mean is more likely to generate low realized returns, and the conditional mean is likely to be persistent. On the other hand, one might think that the expected return has increased, since increases in expected future returns tend to be accompanied by low realized returns. When ρ_{uw} is sufficiently negative, the latter effect outweighs the former and recent returns enter negatively when estimating the current expected return. At the same time, more distant past returns enter positively because they are more informative about the level of the unconditional expected return than about recent changes in the conditional expected return.

We illustrate the role of lagged returns in a simplified setting where historical returns are the only available sample information (D_t). Suppose, for example, that an investor in January 2000 is forming an expectation of the stock market return over the following quarter based on the post-war history of realized market returns. Does the dramatic rise in stock prices in the 1990s increase or decrease the investor's expectation of future return? The answer depends on the extent to which the 1990s' bull market was caused by unexpected declines in expected returns. The conditional expected stock return in this simplified setting is just a weighted average of all past realized returns,

$$\mathrm{E}(r_{t+1} \mid D_t) = \sum_{s=0}^{t-1} \kappa_s r_{t-s},$$

and the weights κ_s depend on ρ_{uw}. For example, if this investor believes that $\rho_{uw} = -0.85$, so that 72% of the variance in unexpected returns is due to changes in expected returns (the estimates of Campbell (1991) are in that neighborhood), then returns realized during the most recent decade receive negative weights, while the returns from the previous four decades receive positive weights. The investor in this example views the 1990s' bull market as a bearish indicator.

Imperfection in predictors complicates inference about their relations to expected return. We show that if predictors are imperfect, the residuals in the predictive regression of r_{t+1} on x_t are serially correlated. This correlation is often ignored when computing standard errors in predictive regressions. The serial correlation in residuals joins other features of predictive regressions that are already well known to complicate inferences, especially in finite samples, such as persistence in the predictors and correlation between the residuals and innovations in the predictors (e.g., Stambaugh (1999)). Using our alternative framework—the predictive system—we develop a Bayesian approach that allows us to conduct clean finite-sample inference about various properties of the expected return. This approach also allows us to incorporate prior beliefs about ρ_{uw}.

A striking example of the importance of such prior beliefs is provided by regressing post-war U.S. stock market returns on the "bond yield," defined as minus the yield on the 30-year Treasury bond in excess of its most recent 12-month moving average. That variable receives a highly significant positive slope (with a p-value of 0.001) in the predictive regression, but its AR(1) innovations are positively correlated with the residuals in that regression. The latter correlation, opposite in sign to what one would anticipate for ρ_{uw}, suggests that the bond yield is a rather imperfect predictor of stock returns. When judged in a predictive system, the bond yield's importance as a predictor depends heavily on prior beliefs about ρ_{uw}. With noninformative beliefs about ρ_{uw}, the bond yield appears to be a very useful predictor; for example, the posterior mode of its conditional correlation with μ_t is 0.9. However, with a more informative belief that innovations in expected returns are negatively correlated with unexpected returns and explain at least half of their variance, the bond yield's conditional correlation with μ_t drops to 0.2. With the more informative belief, the current

value of the bond yield explains only 3% of the variance of μ_t. Adding lagged unexpected returns allows the system to explain 86% of this variance, and further adding lagged predictor innovations increases the fraction of the explained variance of μ_t to 95%.

Prior beliefs also affect the predictive system's advantage in explanatory power over the predictive regression. In the same bond yield example, with noninformative prior beliefs, the predictive system produces an estimate of μ_t that is 1.4 times more precise than the estimate from the predictive regression. With the more informative beliefs, though, the system's estimate is 12.5 times more precise. We measure improvements in precision by improvements in explanatory power. Specifically, we compute the posterior mean of the ratio of the R^2 from the predictive regression to the R^2 from the regression of r_{t+1} on the return forecast from the predictive system.

We also include as predictors two more familiar choices, the market's dividend yield and the consumption-wealth variable "CAY" proposed by Lettau and Ludvigson (2001). Prior beliefs about ρ_{uw} play a less dramatic role with these predictors than with the bond yield, but different prior beliefs can nevertheless produce substantial differences in estimated expected returns. We assess the economic significance of these expected return differences by comparing average certainty equivalents for mean-variance investors whose risk aversion would dictate an all-equity portfolio (i.e., no cash or borrowing) when expected return and volatility equal their long-run sample values. When all three predictors are included, an investor with the more informative belief mentioned above would suffer an average quarterly loss of 1.5% if forced to hold the portfolio selected each quarter by an investor who estimates expected return by the maximum likelihood procedure (which reflects noninformative views about all parameters, including ρ_{uw}).

Ferson, Sarkissian, and Simin (2003) show that persistent predictors may exhibit spurious predictive power in finite samples even if they have no such power in population (e.g., if they have been data-mined). We provide tools that can be helpful in avoiding the spurious regression problem. A spurious predictor is unlikely to produce expected return estimates whose innovations are substantially negatively correlated with unexpected returns. Therefore, under an informative prior about this correlation, a predictive system would likely find the spurious predictor to be almost uncorrelated with μ_t. The basic intuition holds also outside the predictive system framework: If a predictor does not generate a negative correlation between expected and unexpected returns, it is unlikely to be highly correlated with the conditional expected return.

This study is clearly related to an extensive literature on return predictability, but it also contributes to a broader agenda of incorporating economically motivated informative prior beliefs in inference and decision making in finance. Studies in the latter vein include Pástor and Stambaugh (1999, 2000, 2001, 2002a, 2002b), Pástor (2000), Baks, Metrick, and Wachter (2001), and Jones and Shanken (2005). Studies that employ informative priors in the context of return predictability include Kandel and Stambaugh (1996), Avramov (2002,

2004), Cremers (2002), Avramov and Wermers (2006), and Wachter and Warusawitharana (2006).

The paper is organized as follows. Section I introduces the predictive system and discusses its properties. Section II explains why ρ_{uw} is likely to be negative and how it affects expected returns. Section III presents our empirical work, in which we estimate the predictive system with a Bayesian approach. We compare the explanatory powers of the predictive system and predictive regression and quantify the differences in expected return estimates. We also assess the degree to which various predictors are correlated with the expected return. Finally, we decompose the variation in the expected return into three key components. Section IV reviews and concludes.

I. Predictive System

In the predictive regression approach, the expected return is modeled as a linear combination of the predictors in x_t. This modeling assumption is unlikely to be exact, in that no linear combination of the predictors is likely to capture perfectly the unobserved expected return, μ_t. We relax this assumption and develop an alternative predictive framework, which we call a *predictive system*. The predictive regression is a special case of the predictive system, as we show in Section I.C.

We define the predictive system in its most general form as a vector autoregression (VAR) for r_t, x_t, and μ_t, with coefficients restricted so that μ_t is the mean of r_{t+1}. For example, in the first-order VAR, the coefficients relating r_{t+1} to x_t and r_t are set to zero. That predictive system can be viewed alternatively as an unrestricted first-order VAR for r_t, x_t, and a set of unobserved additional predictors, as we explain in the Appendix.[2]

We do not analyze the most general form of the predictive system here. Given our objective to provide an initial exploration of predictive systems, simplicity is a virtue. We examine a simple version of the predictive system, in which μ_t and x_t follow AR(1) processes:

$$r_{t+1} = \mu_t + u_{t+1} \tag{3}$$

$$x_{t+1} = (I - A)E_x + Ax_t + v_{t+1} \tag{4}$$

$$\mu_{t+1} = (1 - \beta)E_r + \beta\mu_t + w_{t+1}. \tag{5}$$

The residuals in the system are assumed to be distributed identically and independently across t as

---

[2] That VAR includes as a special case a first-order VAR for just r_t and x_t, which is a model often employed in research on return predictability. Kandel and Stambaugh (1987) and Campbell (1991) are early examples.

$$\begin{bmatrix} u_t \\ v_t \\ w_t \end{bmatrix} \sim N \left(\begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} \sigma_u^2 & \sigma_{uv} & \sigma_{uw} \\ \sigma_{vu} & \Sigma_{vv} & \sigma_{vw} \\ \sigma_{wu} & \sigma_{wv} & \sigma_w^2 \end{bmatrix} \right). \tag{6}$$

We assume that $0 < \beta < 1$ and that the eigenvalues of A lie inside the unit circle. Denote the covariance matrix in (6) as Σ. Keep in mind that the r_t's and x_t's are observable but the μ_t's are not.

The predictive system is a version of a state space model.[3] Equation (3) defines the unobserved conditional expected return μ_t. Equation (4) is a standard assumption in the predictability literature. A special case of the predictive system arises when there are no predictors, in which case equation (4) is absent and the data include only returns. Equation (5) postulates a simple persistent process for μ_t. This reduced-form model could be consistent with a variety of economic models, rational or behavioral, in which the expected return varies over time in a persistent fashion.

A. Conditional Expected Return

The conditional expected return μ_t is conditioned on a set of information at time t, $\mathcal{F}_t$, that includes at least D_t, the history of returns and predictors observed through time t. The value of μ_t is generally unobservable, even with complete knowledge of the predictive system's parameters—E_r, E_x, A, β, and Σ. Those parameters do, however, imply a value for $\mathrm{E}(\mu_t \,|\, D_t) = \mathrm{E}(r_{t+1} \,|\, D_t)$. Using the Kalman filter, we find that this conditional expected return can be written as the unconditional expected return plus linear combinations of past return forecast errors and innovations in the predictors. Specifically, if we define the forecast error for the return in each period t as

$$\epsilon_t = r_t - \mathrm{E}(r_t \,|\, D_{t-1}), \tag{7}$$

then the expected return conditional on the history of returns and predictors is given by

$$\mathrm{E}(r_{t+1} \,|\, D_t) = E_r + \sum_{s=0}^{\infty} (\lambda_s \epsilon_{t-s} + \phi_s' v_{t-s}), \tag{8}$$

where E_r denotes the unconditional mean return and, in steady state,

$$\lambda_s = m\beta^s \tag{9}$$

$$\phi_s = n\beta^s, \tag{10}$$

---

[3] Harvey (1989) provides a textbook treatment of state-space models, including a brief discussion of the case with nonzero correlations among all of the model's disturbances, which is the case here. In the Appendix, we provide an independent treatment specific to the system in (3) to (6). Studies that analyze return predictability using state-space models include Conrad and Kaul (1988), Lamoureux and Zhou (1996), Johannes, Polson, and Stroud (2002), Ang and Piazzesi (2003), Brandt and Kang (2004), Dangl and Halling (2006), Duffee (2006), and Rytchkov (2007).

where m and n are functions of the parameters in equations (3) to (6).[4] The conditional expected return thus depends on the full history of returns and predictor realizations. We analyze this dependence in more detail in Section II.B, where we plot λ_s as a function of s and ρ_{uw}.

Since the forecast errors ϵ_t in equation (8) are defined relative to conditional expectations that are updated through time based on the available return histories, part of the effects of past return realizations are impounded in those earlier conditional expectations. To isolate the full effect of each past period's total return, we can subtract the unconditional mean from each return, defining $\epsilon_t^U = r_t - E_r$, and then rewrite the conditional expected return in equation (8) as

$$\mathrm{E}(r_{t+1}\,|\,D_t) = E_r + \sum_{s=0}^{\infty} \left(\omega_s \epsilon_{t-s}^U + \delta_s' v_{t-s} \right), \tag{11}$$

where, again in steady state,

$$\omega_s = m(\beta - m)^s \tag{12}$$

$$\delta_s = n(\beta - m)^s. \tag{13}$$

It can be verified that the rate of decay in ω_s and δ_s, $\beta - m$, is nonnegative. This alternative representation of the conditional expected return will also be useful in Section II.B.

B. Temporal Dependence in Returns

Returns in the predictive system exhibit interesting temporal dependence. Given the AR(1) process for μ_t in equation (5), we can rewrite μ_t as an MA(∞) process (the Wold representation):

$$\mu_t = E_r + \sum_{i=0}^{\infty} \beta^i w_{t-i}. \tag{14}$$

Using equations (3) and (14), the return k periods ahead can be written as

$$r_{t+k} = E_r + \sum_{i=0}^{\infty} \beta^i w_{t+k-1-i} + u_{t+k} \tag{15}$$

$$= (1 - \beta^{k-1})E_r + \beta^{k-1}\mu_t + \sum_{i=1}^{k-1} \beta^{k-1-i} w_{t+i} + u_{t+k}. \tag{16}$$

---

[4] In general, m and n are also functions of time, but as the length of the history in D_t grows long, they converge to steady-state values that do not depend on t. That convergence is reached fairly quickly in the settings we consider. We first present the steady-state expressions, for simplicity, but later employ the finite-sample Kalman filter as well. The Appendix derives the functions m and n in finite samples as well as in steady state. The Appendix also shows (in equation A20) that the finite-sample versions of m and n can be interpreted as the slope coefficients from the regression of μ_t on r_t and x_t, respectively, conditional on the sample information at time $t-1$.

This equation implies that the autocovariance of returns is equal to

$$\text{Cov}(r_t, r_{t-k}) = \beta^{k-1}(\beta\sigma_\mu^2 + \sigma_{uw}), \tag{17}$$

where $\sigma_\mu^2 = \sigma_w^2/(1-\beta^2)$ is the unconditional variance of μ_t. As a result, the serial correlation in returns can be positive or negative, depending on the parameter values. The positive component of (17) is due to persistence in μ_t; the negative component is due to $\rho_{uw} < 0$. The knife-edge case of zero autocorrelation obtains for $\rho_{uw} = -\beta\sigma_w/(\sigma_u(1-\beta^2))$.

The AR(1) process for μ_t also implies that returns follow an ARMA(1,1) process,

$$r_{t+1} = (1-\beta)E_r + \beta r_t + \epsilon_{t+1}^* - \gamma\epsilon_t^*. \tag{18}$$

When we implement the Kalman filter without using any predictor information, so that D_t includes only the return history, we obtain equation (8) without the last term ($\sum_{s=0}^{\infty} \phi_s' v_{t-s}$). That equation implies a specification of (18) with $\epsilon_t^* = \epsilon_t$ and $\gamma = \beta - m$.

C. Predictive Regression

The traditional approach to modeling return predictability, a predictive regression

$$r_{t+1} = a + b'x_t + e_{t+1}, \tag{19}$$

arises as a special case of the predictive system if the predictors in x_t are "perfect" in that $\mu_t = a + b'x_t$. The predictors are perfect if there exists a b such that $w_t = b'v_t$ and $A'b = \beta b$.[5] For example, if x_t contains one predictor, this predictor is perfect if its innovations are perfectly correlated with the innovations in μ_t (i.e., $\rho_{vw} = \pm 1$) and if its autocorrelation is the same as that of μ_t (i.e., $A = \beta$). In general, though, the predictors in x_t are imperfect in that $\mu_t \neq a + b'x_t$.

When the predictors approach perfection, then $m \to 0, n \to b$, and equation (8) becomes

$$\text{E}(r_{t+1} \mid D_t) = E_r + b' \sum_{s=0}^{\infty} A^s v_{t-s} = E_r + b'[x_t - E_x] = a + b'x_t, \tag{20}$$

where E_x is the unconditional mean of x_t. That is, when the predictors approach perfection, the system-based conditional expected return approaches the regression-based conditional mean, $a + b'x_t$. When the predictors are imperfect, however, their entire history enters the conditional expected return, since the weighted sum of their past innovations in equation (8) does not then reduce to a function of just x_t. Moreover, when the predictors are imperfect, the

---

[5] $A'b = \beta b$ means that β is an eigenvalue of A' corresponding to the eigenvector b; one example is $A = \beta I$.

expected return depends also on the full history of returns in addition to the history of the predictors.

Predictive systems have interesting implications for predictive regressions. All parameters of the predictive regression in equation (19) can be computed from the parameters of the predictive system in equations (3) to (6).[6] Most interesting, the residual autocovariance is given by

$$\text{Cov}(e_t, e_{t+1}) = \beta\big(\sigma_\mu^2 - V'_{x\mu}V_{xx}^{-1}V_{x\mu}\big) + \sigma_{uw} - V'_{x\mu}V_{xx}^{-1}\sigma_{vu}$$

$$= \beta\text{Var}(\mu_t \,|\, x_t) + \text{Cov}(u_t, w_t - b'v_t). \tag{21}$$

If the predictors are perfect, $\text{Var}(\mu_t \,|\, x_t) = 0$ and $w_t = b'v_t$, so $\text{Cov}(e_t, e_{t+1})$ is zero. With imperfect predictors, though, $\text{Var}(\mu_t \,|\, x_t) > 0$, $w_t \neq b'v_t$, and $\text{Cov}(e_t, e_{t+1})$ is generally nonzero. This observation suggests a simple diagnostic for predictor imperfection: If the residuals from the predictive regression exhibit nonzero autocorrelation, then the predictors x_t are imperfect.

The serial correlation in the residuals complicates the calculation of standard errors in the predictive regression approach. Since the first term in equation (21) is nonnegative, $\text{Cov}(e_t, e_{t+1})$ is often positive, in which case assuming uncorrelated predictive regression residuals leads to understated standard errors. Ferson, Sarkissian, and Simin (2003) make a similar point when the predictor is "spurious," or uncorrelated with expected return. Their setting is a special case of (3) to (6) with one predictor and a diagonal covariance matrix in (6).[7] In specifying a diagonal covariance matrix for the disturbances, they assume not only that the predictor is spurious but also that the innovations in expected return are uncorrelated with unexpected returns (i.e., $\rho_{uw} = 0$). In this special case, we see from (21) that $\text{Cov}(e_t, e_{t+1}) = \beta\sigma_\mu^2$. Ferson et al. do not report this expression but do find, using simulations, that the positive residual serial correlation can substantially affect inference in predictive regressions. Duffee (2006) also uses simulations to make a related point in the context of bond predictability.

II. Correlation between Expected and Unexpected Returns

A key quantity in this paper is ρ_{uw}, the correlation between the unexpected return, u_t, and the innovation in the expected return, w_t. Henceforth, we refer to ρ_{uw} simply as the "correlation between expected and unexpected returns," a slightly inaccurate but much shorter description. This correlation is important for return predictability in several ways. First, it determines how past returns affect the forecasts of future returns, as we show in Section II.B. Second, prior beliefs about ρ_{uw} play an important role in various inferences about predictability, as we show in Section III. The ability to incorporate prior beliefs about ρ_{uw}

---

[6] For example, the regression slope b can be computed from the system's parameters as $b = V_{xx}^{-1}V_{x\mu}$, where V_{xx} is given in the Appendix in equation (A33), $V_{x\mu} = (I_K - \beta A)^{-1}\sigma_{vw}$, and I_K is a $K \times K$ identity matrix.

[7] The objectives of Ferson et al. differ from ours. For example, they do not use this multiple-equation setting to estimate expected return or to examine its dependence on lagged returns and predictors.

is a key feature of the predictive system. We begin this section by discussing some theoretical properties of ρ_{uw}.

A. Why ρ_{uw} Is Likely to Be Negative

The basic motivation behind the belief that $\rho_{uw} < 0$ is that asset prices tend to fall when discount rates rise. More precisely, $\rho_{uw} < 0$ means that unanticipated increases in expected returns tend to be accompanied by unexpected negative returns. This intuition holds perfectly for nominal returns on Treasury bonds. Since the nominal cash flows of Treasury bonds are fixed, the bond price variation is driven only by discount rate shocks, and $\rho_{uw} = -1$. Stock returns, however, are driven also by cash flow shocks. It is possible, at least in principle, that positive discount rate shocks could be accompanied by such large positive cash flow shocks that stock prices rise rather than fall as a result. In this section, we derive conditions under which $\rho_{uw} < 0$, and argue that these conditions are likely to be satisfied for the aggregate stock market.

Following Campbell (1991), the unexpected return can be decomposed approximately as

$$u_{t+1} = \eta_{C,t+1} - \eta_{E,t+1}, \tag{22}$$

where $\eta_{C,t+1}$ represents the unanticipated revisions in expected future cash flows and $\eta_{E,t+1}$ captures the revisions in expected future returns. If the expected return follows the process in equation (2) with $0 < \beta < 1$, then $\eta_{E,t+1}$ is equal to w_{t+1} multiplied by a positive constant, so that

$$\rho_{uw} = \rho(u_{t+1}, \eta_{E,t+1}). \tag{23}$$

Since the partial correlation between u_{t+1} and $\eta_{E,t+1}$ is minus one (equation (22)), it seems natural to believe a priori that the simple correlation between u_{t+1} and $\eta_{E,t+1}$, or ρ_{uw}, is negative. Before seeing any data, it is not obvious why $\eta_{C,t+1}$ and $\eta_{E,t+1}$ should be correlated, and a belief that $\eta_{C,t+1}$ and $\eta_{E,t+1}$ are uncorrelated translates into a belief that ρ_{uw} is negative. More precisely, it follows directly from equations (22) and (23) that $\rho_{uw} < 0$ if and only if

$$\rho(\eta_{C,t+1}, \eta_{E,t+1}) < \frac{\sigma(\eta_{E,t+1})}{\sigma(\eta_{C,t+1})}, \tag{24}$$

where the σ's denote standard deviations (*SD*). In order for $\rho_{uw} < 0$ to be violated, cash flow shocks would have to be more important than discount rate shocks in explaining the variance of stock returns, that is, $\sigma(\eta_{C,t+1}) > \sigma(\eta_{E,t+1})$, and the correlation between those shocks, $\rho(\eta_{C,t+1}, \eta_{E,t+1})$, would have to be positive and sufficiently high. It seems difficult to argue that one could expect such a high correlation a priori.[8] In fact, such a high correlation between the

---

[8] Consistent with our perspective, Campbell (2001) argues that negative ρ_{uw} "is of special interest." Campbell's table 1 reports the implied values of ρ_{uw} for various combinations of parameters related to the same Campbell (1991) decomposition that we use in our equation (22). All values of ρ_{uw} reported in Campbell's table 1 are negative.

shocks to cash flows and discount rates seems unlikely because it would make stock returns unrealistically smooth. It is easy to see from equation (22) that a violation of the condition in (24) would require that

$$\text{Var}(u_{t+1}) < \text{Var}(\eta_{C,t+1}) - \text{Cov}(\eta_{C,t+1}, \eta_{E,t+1})$$
$$< \text{Var}(\eta_{C,t+1}) - \text{Var}(\eta_{E,t+1}) < \text{Var}(\eta_{C,t+1}). \tag{25}$$

That is, for $\rho_{uw} < 0$ to be violated, stock returns would have to be less volatile than when the expected return is constant (i.e., when $\text{Var}(\eta_{E,t+1}) = 0$).

In reality, though, stock returns appear to be more volatile than when the expected return is constant. For example, Shiller (1981) finds that stock returns are much more volatile than the present value of future dividends discounted at constant rates. To explain this "excess volatility puzzle," discount rates must vary over time in a way that increases stock volatility. But if ρ_{uw} were positive, discount rate variation would actually reduce stock volatility, thereby deepening the puzzle.[9] Shiller's results largely pre-date the sample period that we use in our empirical work in Section III—only one quarter of Shiller's sample period (1871 to 1979) overlaps with ours, which begins in 1952. Therefore, we feel comfortable using Shiller's results as one source of prior information suggesting $\rho_{uw} < 0$.

Another source of prior information suggesting $\rho_{uw} < 0$ is the evidence of Campbell (1991). Campbell uses a vector-autoregressive approach to decompose unexpected stock market returns into components due to cash flow shocks, $\eta_{C,t+1}$, and discount rate shocks, $\eta_{E,t+1}$, as in our equation (22). Campbell considers two subperiods, 1927 to 1951 and 1952 to 1988, so we can view his results from the earlier subperiod as prior information for our empirical analysis beginning in 1952. Based on quarterly data (which we also use in our empirical work), Campbell estimates in his Table 2 that $\sigma(\eta_{E,t+1}) > \sigma(\eta_{C,t+1})$ in the period 1927 to 1951, meaning that discount rate news is more important than cash flow news in explaining the variance of stock market returns. This result makes the condition (24) hold trivially, since $\rho(\eta_{C,t+1}, \eta_{E,t+1}) < 1$. Moreover, Campbell obtains negative estimates of $\rho(\eta_{C,t+1}, \eta_{E,t+1})$ in the period 1927 to 1951, which again makes (24) hold trivially independent of $\sigma(\eta_{E,t+1})/\sigma(\eta_{C,t+1})$. In fact, Table 2 of Campbell (1991) points to large negative estimates of ρ_{uw}. He reports estimates of the variance of $\eta_{E,t+1}$, $\sigma_{\eta_E}^2$, and its covariance with $\eta_{C,t+1}$, $\sigma(\eta_C, \eta_E)$, both as fractions of σ_u^2, from which the implied estimates of $\rho(u_{t+1}, \eta_{E,t+1})$ can be computed as $[\sigma(\eta_C, \eta_E)/\sigma_u^2 - \sigma_{\eta_E}^2/\sigma_u^2]/[\sigma_{\eta_E}/\sigma_u]$. Given equation (23), we interpret these values as implied estimates of ρ_{uw}. For the period 1927 to 1951,

---

[9] Our argument is not fully precise because Shiller analyzes the unconditional variance of stock returns, whereas $\text{Var}(u_{t+1})$ in equation (25) represents the conditional variance. But the difference between the two variances is relatively small because the variance of μ_t is generally agreed to be much smaller than the variance of u_{t+1} in equation (1). Moreover, we can allow plenty of margin for error on both sides of our argument. On the one side, $\rho_{uw} > 0$ implies $\text{Var}(u_{t+1}) < \text{Var}(\eta_{C,t+1})$ with as many as three inequality signs in equation (25). On the other side, Shiller's evidence that $\text{Var}(u_{t+1}) > \text{Var}(\eta_{C,t+1})$ seems quite strong (see Shiller's Figure 1).

Campbell's results imply values of ρ_{uw} ranging from -0.67 to -0.87 across three different specifications.

Campbell uses three predictors: the dividend-price ratio (D/P), the lagged stock return, and the (relative) 1-month T-bill rate. He reports that his results are sensitive to the exclusion of D/P but robust as long as D/P is included. To assess further the sensitivity of the pre-1952 evidence about ρ_{uw} to the choice of predictors, we estimate predictive regressions based on quarterly data in 1927 to 1951 with various combinations of five predictors: D/P, the AAA-BAA default spread, the term spread, and both the long- and short-term government bond yields in excess of their 12-month moving averages. We find large negative estimates of ρ_{uw} in all cases when D/P is included and in most cases when it is not. Tabulated results are omitted to save space.

The empirical evidence from earlier periods, coupled with the basic discount rate motivation, suggests that $\rho_{uw} < 0$ is a sensible prior belief for our empirical work. To broaden our investigation, however, we entertain three different priors on ρ_{uw}, including a noninformative prior.

Finally, we should note that various studies provide post-war empirical evidence that $\rho_{uw} < 0$. The implied estimates of ρ_{uw} from Table 2 of Campbell (1991), discussed earlier, range from -0.92 to -0.94 in the later subperiod (1952 to 1988) and from -0.71 to -0.86 in the full sample (1927–1988). Campbell and Ammer (1993) use seven predictors and report estimates in their Table III that imply estimates of ρ_{uw} ranging from -0.93 to -0.95. van Binsbergen and Koijen (2007) estimate ρ_{uw} ranging from -0.67 to -0.45. They do not rely on pre-specified predictors but instead use data on dividend growth and returns in the context of a present value model. They also find a positive correlation between shocks to expected return and dividend growth, similar to Menzly, Santos, and Veronesi (2004), Lettau and Ludvigson (2005), and Kothari, Lewellen, and Warner (2006). Note that $\rho(\eta_{C,t+1}, \eta_{E,t+1})$ can be positive without violating the condition in equation (24).

B. *The Role of ρ_{uw} in Determining Expected Returns*

The correlation between expected and unexpected returns, ρ_{uw}, plays a critical role in determining the conditional expected return. To illustrate this role, we consider a special case of the predictive system in which there are no predictors. With no predictors, D_t includes only the return history, and the conditional expected return in equation (8) is simply $\mathrm{E}(r_{t+1} \mid D_t) = E_r + \sum_{s=0}^{\infty} \lambda_s \epsilon_{t-s}$, a weighted sum of past forecast errors in returns (the Wold representation). Panel A of Figure 1 plots the values of λ_s in an example with the predictive R^2 (the fraction of the variance in r_{t+1} explained by μ_t) equal to 0.05, β equal to 0.9, and four different values of ρ_{uw} ranging from -0.99 to 0. The figure shows that different values of ρ_{uw} produce different values of m, and hence also different behaviors for $\lambda_s (= m\beta^s)$.

The results in Figure 1 can be understood by noting that there are essentially two effects of the return history on the current expected return. The first might be termed the "level" effect. Observing recent realized returns that were higher than expected suggests that they were generated from a distribution with a
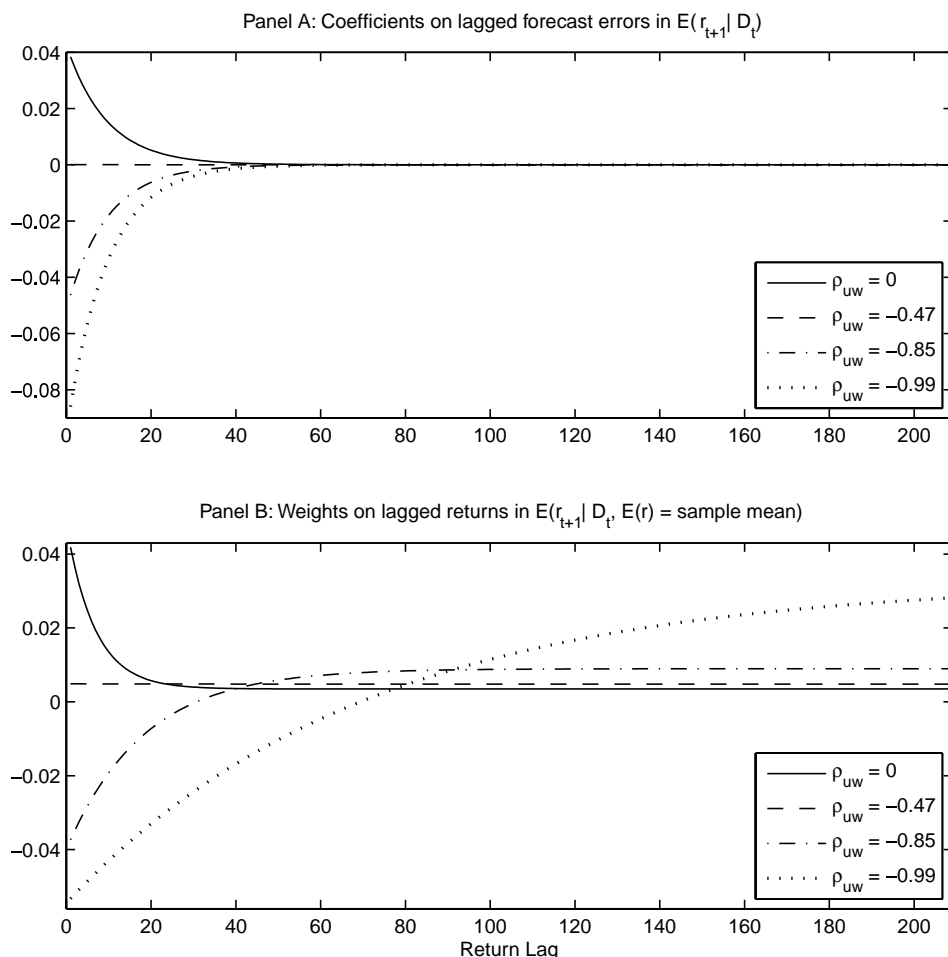
Panel A: Coefficients on lagged forecast errors in $E(r_{t+1}| D_t)$

Panel B: Weights on lagged returns in $E(r_{t+1}| D_t, E(r) = $ sample mean)

Return Lag

Figure 1. The effect of lagged returns on $E(r_{t+1} | D_t)$ when no predictors are used. Panel A plots λ_s, the coefficients on lagged forecast errors ($\epsilon_{t-s} = r_{t-s} - E(r_{t-s} | D_{t-s-1})$) in $E(r_{t+1} | D_t)$, where r_{t+1} denotes the stock return from time t to time $t + 1$ and D_t denotes the history of returns and predictors observed through time t. Panel B plots κ_s, the weights on lagged total returns in $E(r_{t+1} | D_t)$ when the unconditional mean return is estimated by the sample mean over the previous 208 quarters (which is the length of the sample used in subsequent analysis). No predictors are used in the predictive system. The steady-state values of all coefficients are plotted. The different lines correspond to different values of ρ_{uw}, the correlation between expected and unexpected returns. The mean reversion coefficient in the AR(1) process for the conditional expected return μ_t is set equal to $\beta = 0.9$. The predictive R^2—the fraction of variation in r_{t+1} that can be explained by μ_t—is set equal to $R^2 = 0.05$.

higher mean. If the expected return is persistent, as it is in this example with $\beta = 0.9$, then that recent history suggests that the current mean is higher as well. So the level effect positively associates past forecast errors in returns with expected future returns. The second effect, which might be termed the "change" effect, operates via the correlation between expected and unexpected returns. In

particular, suppose ρ_{uw} is negative, as we suggest is reasonable. Then observing recent realized returns that were higher than expected suggests that expected returns fell in those periods. That is, part of the reason that realized returns were higher than expected is that there were price increases associated with negative shocks to expected future returns and thus to discount rates applied to expected future cash flows. So the change effect negatively associates past forecast errors in returns with expected future returns. Overall, the net impact of the return history on the current return depends on the relative strengths of the level and change effects.

The level and change effects can be mapped into the return autocovariance in (17). When ρ_{uw} is sufficiently negative, then $\beta\sigma_\mu^2 < -\sigma_{uw}$, returns are negatively autocorrelated, and the change effect prevails. Also, $m < 0$ in that case, so the λ_s's in (9) are negative. When $\beta\sigma_\mu^2 > -\sigma_{uw}$, returns are positively autocorrelated, the λ_s's are positive, and the level effect prevails.

When $\rho_{uw} = 0$, there is no change effect and only the level effect is present. For that case, the λ_s's in Figure 1 start at a positive value for the first lag, about 0.04, and then decay toward zero. The level and change effects offset each other when $\rho_{uw} = -0.47$ (this is the knife-edge case of zero autocorrelation in equation (17)), or when the fraction of the variance in unexpected returns explained by expected return shocks, ρ_{uw}^2, is about 22%. In that case, the λ_s's plot as a flat line at zero. This result is worth emphasizing: For $\rho_{uw} = -0.47$, rational investors who know the unconditional expected return do not update their beliefs about the conditional expected return, regardless of what realized returns they observe. The change effect dominates when $\rho_{uw} = -0.85$, where the λ_s's start around -0.04 at the first lag, and it is even stronger when $\rho_{uw} = -0.99$, where the λ_s's start around -0.08. Clearly, the correlation between expected and unexpected returns is a critical determinant of the relation between the return history and the current expected return.

The conditional expected return depends on the true unconditional mean, E_r, which must be estimated in practice. A natural estimator is the sample mean. Consider again the no-predictor case where the summation on the right-hand side of equation (11) is truncated at $s = t - 1$ and E_r is replaced by the sample mean, $(1/t)\sum_{l=1}^{t} r_l$. The estimated conditional expected return then becomes a weighted average of past returns,

$$E(r_{t+1} \mid D_t) = \sum_{s=0}^{t-1} \kappa_s r_{t-s}, \tag{26}$$

where

$$\kappa_s = \frac{1}{t}\left(1 - \sum_{l=1}^{t} \omega_l\right) + \omega_s \tag{27}$$

and $\sum_{s=0}^{t-1} \kappa_s = 1$. The weights ($\kappa_s$'s) are plotted in Panel B of Figure 1 for $t = 208$, corresponding to the number of quarters used in our empirical analysis. When $\rho_{uw} = 0$, all past returns enter positively but recent returns are weighted more heavily. In the $\rho_{uw} = -0.47$ case, where the level and change effects exactly

offset each other, all of the weights equal $1/t$, so the conditional expected return is then just the historical sample average. For the larger negative ρ_{uw} values, where the change effect is stronger, the weights switch from negative at more recent lags to positive at more distant lags (as the weights must sum to one). For example, when changes in expected returns explain about 72% of the variance in unexpected returns ($\rho_{uw} = -0.85$), the returns from the most recent 10 years (40 quarters) contribute negatively to the estimated current expected return, while the returns from the earlier 42 years contribute positively.

An additional perspective on the role of ρ_{uw} is provided by the time series of conditional expected returns plotted in Figure 2. In constructing these series,

Figure 2. The equity premium $\mathbf{E}(r_{t+1} \,|\, D_t)$ from the predictive system with no predictors. This figure plots the time series of the quarterly equity premium estimated for four different values of ρ_{uw}, the correlation between expected and unexpected returns. The mean reversion coefficient in the AR(1) process for the conditional expected return μ_t is set equal to $\beta = 0.9$. The predictive R^2—the fraction of variation in r_{t+1} that can be explained by μ_t—is set equal to $R^2 = 0.05$. The parameters represent quarterly values.

we maintain the same setting and same parameter values as in Figure 1. The unconditional mean return E_r is set equal to the sample average for our 208-quarter sample period, and then, starting from the first quarter in the sample, the conditional mean is updated through time using the finite-sample Kalman filter applied to the realized returns data. As before, the level and change effects exactly offset each other when $\rho_{uw} = -0.47$, so the conditional expected return in that case is simply a flat line at the sample average for the period. A striking feature of the plot is that the expected return series for $\rho_{uw} = 0$ is virtually the mirror image of the series for $\rho_{uw} = -0.85$. Moreover, the differences among the various series of conditional expected returns are large in economic terms, often several percent per quarter. As before, we see that ρ_{uw} plays a key role in estimating expected returns.

Figure 3 compares the R^2s from three approaches to predicting r_{t+1} using some or all of the information observable at time t, which includes the history of returns and a single predictor x_t. The first approach is the predictive system, which uses all of that available history. The second is the predictive regression, which uses only the current value x_t. The third approach is the ARMA(1,1) model in equation (18), which uses only past returns. The four panels correspond to the values {0, 0.3, 0.6, 0.9} for ρ_{vw}, the conditional correlation between μ_t and the single predictor x_t. In all four panels, $\beta = A = 0.9$, and the true predictive R^2 (from the regression of r_{t+1} on μ_t) is 0.05. We consider two values of ρ_{uv}, "high" and "low," which correspond to partial correlations between u_t and v_t given w_t of $\rho_{uv\,|\,w} = 0.9$ and $\rho_{uv\,|\,w} = -0.9$, respectively.[10]

Since all three approaches compared in Figure 3 use only information observable at time t, they all produce R^2s smaller than 0.05. The R^2 from the predictive regression rises from zero to 0.04 as ρ_{vw} rises from zero to 0.9 across the four panels. This increase is intuitive: As μ_t and x_t become more highly correlated, the predictive regression becomes more useful in predicting returns. The predictive regression R^2 is invariant to ρ_{uw}. In contrast, the R^2 from the ARMA(1,1) model, which summarizes the usefulness of past returns in predicting future returns, is heavily influenced by ρ_{uw}. When $\rho_{uw} = -0.47$, this R^2 is zero: Past returns contain no information about future returns because the level and change effects cancel out. For $\rho_{uw} \neq -0.47$, stock returns are serially correlated and the ARMA(1,1) R^2 is positive; in fact, it can be higher than the predictive regression R^2. For example, when $\rho_{vw} = 0.3$ and $\rho_{uw} \notin (-0.74, -0.13)$, past returns are more useful than x_t in predicting r_{t+1}. The highest R^2's are invariably achieved by the predictive system, which uses more information to predict future returns than do the other two approaches.

III. Empirical Analysis

In this section, we use the predictive system to conduct an empirical analysis of return predictability. We first present evidence from predictive regressions,

---

[10] We specify the partial correlation $\rho_{uv\,|\,w}$ instead of the simple correlation ρ_{uv} because our control over ρ_{uv} is limited. The permissible range of values for ρ_{uv} depends on ρ_{vw} and ρ_{uw} (see equation (32)) and we vary both ρ_{vw} and ρ_{uw}. By choosing $\rho_{uv\,|\,w}$ of 0.9 and −0.9, we are specifying ρ_{uv} close to the boundaries of its permissible range.
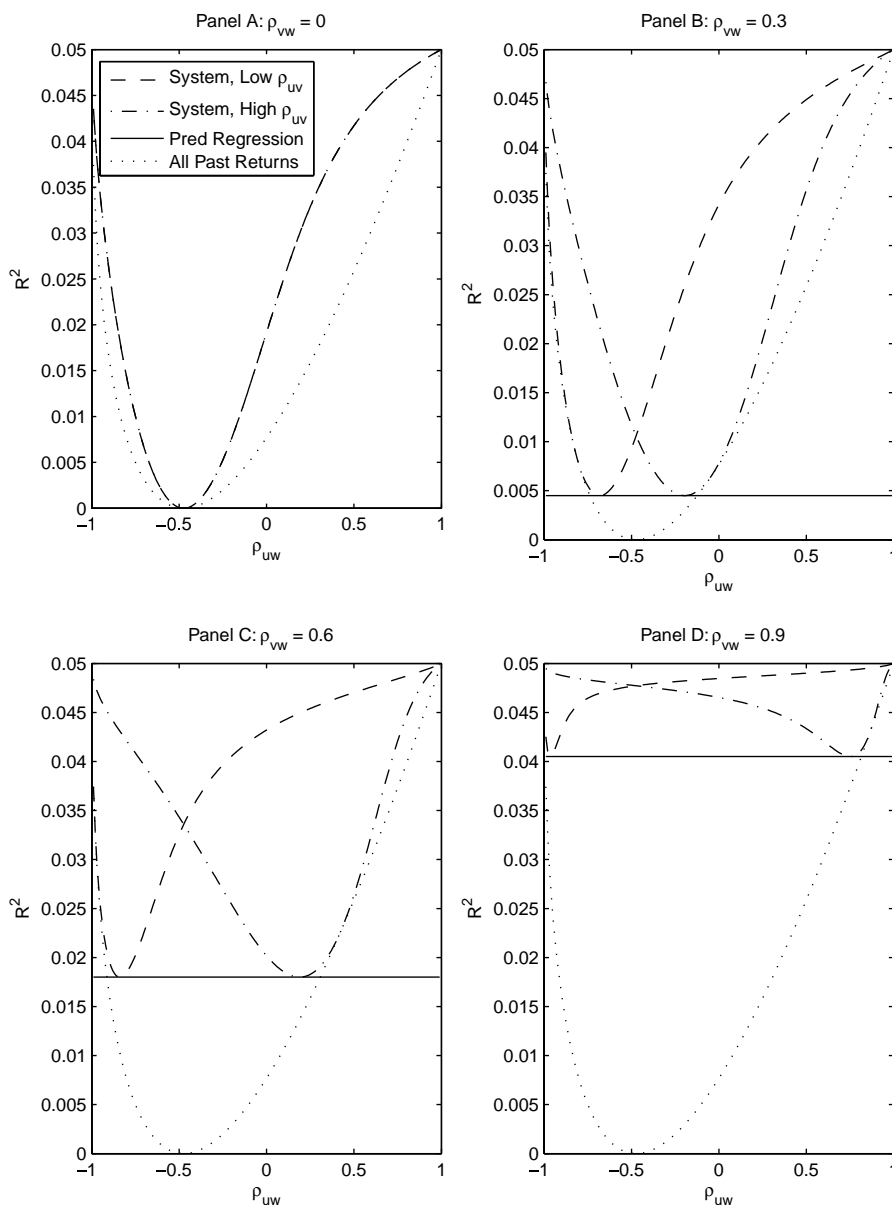
Figure 3. Predictive R^2s. Each panel plots the R^2s from three approaches to predicting stock returns r_{t+1} using information observable at time t. The approaches are: the predictive regression of r_{t+1} on a single predictor x_t (solid line); the ARMA(1,1) model that uses the full history of past returns but no predictor data (dotted line); and the predictive system, which uses the full history of returns and predictor realizations (dashed and dash-dot lines). The dashed (dash-dot) line corresponds to a "low" ("high") value of ρ_{uv}, which represents the value obtained when the partial correlation between u_t and v_t given w_t equals $\rho_{uv\,|\,w} = -0.9\,(0.9)$. The conditional correlation between μ_t and x_t, ρ_{vw}, ranges from zero in Panel A to 0.9 in Panel D. In all four panels, $\beta = A = 0.9$, and the true predictive R^2 (from the regression of r_{t+1} on μ_t) is 0.05. In Panel A, the solid line coincides with the x axis, and the dashed and dash-dot lines overlap.

for benchmark purposes. Then we discuss identification issues and use the system to estimate expected returns via maximum likelihood. Finally, we turn to the main analysis, which takes a Bayesian approach.

A. Evidence from Predictive Regressions

We begin by estimating predictive regressions on quarterly data in 1952 to 2003 for three predictors.[11] The first predictor is the market-wide dividend yield, which is equal to total dividends paid over the previous 12 months divided by the current total market capitalization. We compute the dividend yield from the with-dividend and without-dividend monthly returns on the value-weighted portfolio of all New York Stock Exchange (NYSE), Amex, and Nasdaq stocks, which we obtain from the Center for Research in Security Prices (CRSP) at the University of Chicago. The second predictor is CAY from Lettau and Ludvigson (2001), whose updated quarterly data we obtain from Martin Lettau's web site. The third predictor is the "bond yield," which we define as minus the yield on the 30-year Treasury bond in excess of its most recent 12-month moving average. The bond yield data are from the Fixed Term Indices in the CRSP Monthly Treasury file. The three predictors are used to predict quarterly returns on the value-weighted portfolio of all NYSE, Amex, and Nasdaq stocks in excess of the quarterly return on a 1-month T-bill, which is also obtained from CRSP.

Whereas the first two predictors have been used extensively, the third predictor appears to be new. On the one hand, this predictor has some a priori motivation. It seems plausible for the long-term Treasury bond yield to be related to future stock returns since expected returns on stocks and bonds may comove due to discount rate-related factors. Moreover, subtracting the 12-month average yield is an adjustment that is commonly applied to the short-term risk-free rate (e.g., Campbell (1991) and Campbell and Ammer (1993)). On the other hand, there is no ex ante reason to choose this particular predictor over plausible alternatives such as the term spread, credit spread, etc. We choose the bond yield as the third predictor simply because it illustrates well the effect of prior beliefs about ρ_{uw} on inference about predictability, as shown below.

Table I reports, for various predictive regressions, the estimated slope coefficient vector $\hat{b}$, the R^2, and the estimated correlation between unexpected returns and the innovations in expected returns. This correlation, which represents the regression-based counterpart of ρ_{uw}, is computed as $\text{Corr}(\hat{e}_t, \hat{b}'\hat{v}_t)$, following equations (4) and (19). Table I also reports the ordinary least squares (OLS) t-statistics and the bootstrapped p-values associated with these t-statistics as well as with the R^2.[12]

---

[11] Since we use the T-bond and T-bill yields in our analysis, we begin our sample in 1952, after the 1951 Treasury-Fed accord that made possible the independent conduct of monetary policy. Campbell and Ammer (1993), Campbell and Yogo (2006), and others also begin their samples in 1952 for this reason.

[12] Our bootstrap repeats the following procedure 20,000 times: (i) resample T pairs of $(\hat{v}_t, \hat{e}_t)$, with replacement, from the OLS residuals in regressions (4) and (19); (ii) build up the time series of x_t, starting from the unconditional mean and iterating forward on equation (4), using the OLS estimate $\hat{A}$, sample mean $\hat{E}_x$, and the resampled values of $\hat{v}_t$; (iii) construct the time series of

<div align="center">

Table I

Predictive Regressions
</div>

This table summarizes the results from predictive regressions, $r_t = a + b'x_{t-1} + e_t$, where r_t is the quarterly excess stock market return and x_{t-1} contains the predictors (listed in the column headings) lagged by one quarter. The sample period covers 1952 (Q1) to 2003 (Q4). The table reports the estimated slope coefficients $\hat{b}$, the (unadjusted) R^2 of the predictive regression, and the correlation between the estimated unexpected return $\hat{e}_t$ and the estimated shock to expected return $\hat{b}'\hat{v}_t$. The value of $\hat{v}_t$ is the fitted residual vector from the vector autoregression, $x_t = (I - A)E_x + Ax_{t-1} + v_t$. The R^2 and correlation are reported in percent (i.e., $\times 100$). The OLS t-statistics are given in parentheses "()." The t-statistic of $\text{Corr}(\hat{e}_t, \hat{b}'\hat{v}_t)$ is computed as the t-statistic of the slope from the regression of the sample residuals $\hat{e}_t$ on $\hat{b}\hat{v}_t$. The p-values associated with all t-statistics and R^2s are computed by bootstrapping and reported in brackets "[ ]."

Bond Yield	Dividend Yield	CAY	$\text{Corr}(\hat{e}_t, \hat{b}'\hat{v}_t)$	R^2
2.716			21.735	4.231
(3.024)			(3.204)	[0.002]
[0.001]			[0.001]	
	1.153		−91.887	2.252
	(2.184)		(−33.506)	[0.059]
	[0.057]		[1.000]	
		1.704	−53.556	7.292
		(4.035)	(−9.124)	[0.000]
		[0.000]	[1.000]	
2.573	1.028	1.346	−35.635	11.777
(2.902)	(1.966)	(3.139)	(−5.487)	[0.000]
[0.003]	[0.058]	[0.003]	[1.000]	

The results suggest that all three predictors have some forecasting ability. The dividend yield is marginally significant, whereas both the bond yield and CAY are highly significant predictors. When used alone, both of the latter predictors exhibit p-values of 0.1% or less, and they are about equally significant when all three predictors are used together. It is also informative to examine the correlation between estimated unexpected and expected return, $\text{Corr}(\hat{e}_t, \hat{b}'\hat{v}_t)$, shown in the fourth column of Table I. When the single predictor is either the dividend yield or CAY, this correlation is negative and highly significant: -91.9% for the dividend yield and -53.6% for CAY. These negative correlations are not surprising since both predictors are negatively related to stock prices, by construction. For the bond yield, however, this correlation is positive (21.7%) and highly significant. This positive correlation makes it unlikely that the bond yield is perfectly correlated with the conditional expected return μ_t.

The correlation between expected and unexpected returns is a useful diagnostic that should be considered when examining the output of a predictive

returns, r_t, by adding the resampled values of $\hat{e}_t$ to the sample mean (i.e., under the null that returns are not predictable); and (iv) use the resulting series of x_t and r_t to estimate regressions (4) and (19) by OLS. The bootstrapped p-value associated with the reported t-statistic (or R^2) is the relative frequency with which the reported quantity is smaller than its 20,000 counterparts bootstrapped under the null of no predictability.

regression. Since this correlation is likely to be negative, predictive models in which this correlation is positive seem less plausible.[13] The model in which the bond yield is the single predictor is a good example. Based on the predictive-regression p-value, the bond yield would appear to be a highly successful predictor whose forecasting ability is better than that of the dividend yield and comparable to that of CAY. However, the bond yield produces expected return estimates whose innovations are positively correlated with unexpected returns, suggesting that this predictor is imperfect. We suspect that the same statement can be made about many macroeconomic variables that the literature has related to expected returns. In the rest of the paper, we develop a predictive framework that allows us to incorporate the prior belief that the correlation between expected and unexpected returns is negative.

B. Identification and Maximum Likelihood Estimation

In the absence of any priors or parameter restrictions, not all of the parameters in equations (3) to (6) are identified. We can nevertheless obtain estimates of conditional expected returns using equation (4) and the recursive representation for returns,

$$r_{t+1} = (1 - \beta)E_r + \beta r_t + n'v_t - (\beta - m)\epsilon_t + \epsilon_{t+1}, \tag{28}$$

which follows directly from the steady-state representation of the conditional expected return in (8). The parameters in (4) and (28) are identified and can be estimated using maximum likelihood by representing the two equations as a state-space system and applying standard methodology (e.g., Hamilton (1994), Section 13.4).[14] The parameters in these equations, along with the covariance matrix of $[\epsilon_t \, v_t']$, identify the parameters appearing in equations (3) to (5) but not all of the parameters in the covariance matrix in (6). Only Σ_{vv} is identified just by the data. Identifying the remaining elements of Σ requires additional information about at least one of them.[15]

Figure 4 plots the time series of expected returns obtained via maximum likelihood estimation as well as the expected return estimates obtained from OLS estimation of the predictive regression. Panels A and B display results with a single predictor, either the dividend yield or CAY. In Panel C, these variables are combined with the bond yield variable in the three-predictor case. First, observe that the fluctuation of the expected return estimates seems too large to be plausible. In Panel B, for example, expected returns range from -5% to 8% per quarter, and the range is even wider in Panel C. Later on, we obtain smoother time series of μ_t by specifying informative prior beliefs.

---

[13] Strictly speaking, the arguments based on equations (22) and (24) apply when r_{t+1} denotes the total stock return, but they should hold to a close approximation also when r_{t+1} denotes the excess stock return, as used here. For excess returns, Campbell (1991) shows that equation (22) has an additional term representing news about future interest rates, and he estimates the variance of that term to be an order of magnitude smaller than the variances of $\eta_{C,t+1}$ and $\eta_{E,t+1}$.

[14] The likelihood function is detailed in the online Internet Appendix in the "Supplements and Datasets" section at http://www.afajof.org/supplements.asp.

[15] Rytchkov (2007) discusses identification issues in a similar setting.

Figure 4. The equity premium: Regression versus system with no prior information.
This figure plots the time series of the quarterly equity premium estimated in two different environments. The dotted line plots the OLS fitted values from the predictive regression of r_{t+1} on the given predictor(s). The solid line plots the maximum likelihood estimates of $\mathrm{E}(r_{t+1}\,|\,D_t)$ from the predictive system. In Panel A, the estimation uses one predictor, dividend yield. In Panel B, the single predictor is CAY. In Panel C, three predictors are used: dividend yield, CAY, and the bond yield. The sample period is 1952Q1 to 2003Q4.

Second, observe that although the series of estimated expected returns exhibit marked differences across the three sets of predictors, the differences between the predictive regression estimates and the predictive system estimates for a given set of predictors are much smaller.[16]

---

[16] We also estimate expected returns from the predictive system under diffuse priors (the discussion of prior beliefs follows later in the text). We find that the resulting estimates (not plotted here) behave similarly to both the OLS estimates from the predictive regression and the maximum likelihood estimates from the predictive system.

As the length of the sample grows, posterior beliefs about the parameters in (28) and thus (8) will converge to values that do not depend on prior beliefs about the parameters in the predictive system (as long as those priors do not strictly preclude such values). Therefore, after observing a sufficiently long sample, prior beliefs about ρ_{uw}, for example, will not impact forecasts of future returns. (Our actual sample is evidently not long in that sense, as prior beliefs about ρ_{uw} exert a substantial effect on estimates of expected returns.) On the other hand, given the lack of full identification of Σ, prior beliefs about ρ_{uw} will matter even in large samples when making inferences about the correlation between the predictors and the unobservable expected return μ_t.

C. Bayesian Approach

We develop a Bayesian approach for estimating the predictive system. This approach has several advantages over frequentist alternatives such as the maximum likelihood approach. First, the Bayesian approach allows us to specify economically motivated prior distributions for the parameters of interest. Second, it produces posterior distributions that deliver finite-sample inferences about relatively complicated functions of the underlying parameters, such as the correlations between μ_t and x_t and the R^2s from the regression of r_{t+1} on μ_t. Finally, it incorporates parameter uncertainty as well as uncertainty about the path of the unobservable expected return μ_t.

We obtain posterior distributions using *Gibbs sampling*, a Markov Chain Monte Carlo (MCMC) technique (e.g., Casella and George (1992)). In each step of the MCMC chain, we first draw the parameters $(E_x, A, E_r, \beta, \Sigma)$ conditional on the current draw of $\{\mu_t\}$, and then we use the *forward filtering, backward sampling* algorithm developed by Carter and Kohn (1994) and Frühwirth-Schnatter (1994) to draw the time series of $\{\mu_t\}$ conditional on the current draw of $(E_x, A, E_r, \beta, \Sigma)$.

We impose informative prior distributions on three quantities: the correlation ρ_{uw} between expected and unexpected returns, the persistence β of the expected return μ_t, and the predictive R^2 from the regression of r_{t+1} on μ_t. These prior distributions are plotted in Figure 5.

The key prior distribution is the one on ρ_{uw}. We consider three priors on ρ_{uw}, all of which are plotted in Panel A of Figure 5. The "noninformative" prior is flat on most of the $(-1, 1)$ range, with prior mass tailing off near ± 1 to avoid potential singularity problems. The "less informative" prior imposes $\rho_{uw} < 0$ in that 99.9% of the prior mass of ρ_{uw} is below zero. As shown in Panel B, this prior implies a relatively noninformative prior on ρ_{uw}^2, with most prior mass between zero and 0.8. Finally, the "more informative" prior on ρ_{uw} is specified such that the implied prior on ρ_{uw}^2 has 99.9% of its mass above 0.5, with a mean of about 0.77. Since ρ_{uw}^2 is the R^2 from the regression of unexpected returns on shocks to expected returns, it represents the fraction of market variance that is due to news about discount rates. Therefore, the more informative prior reflects the belief that at least half of the variance of market returns is due to discount rate news.
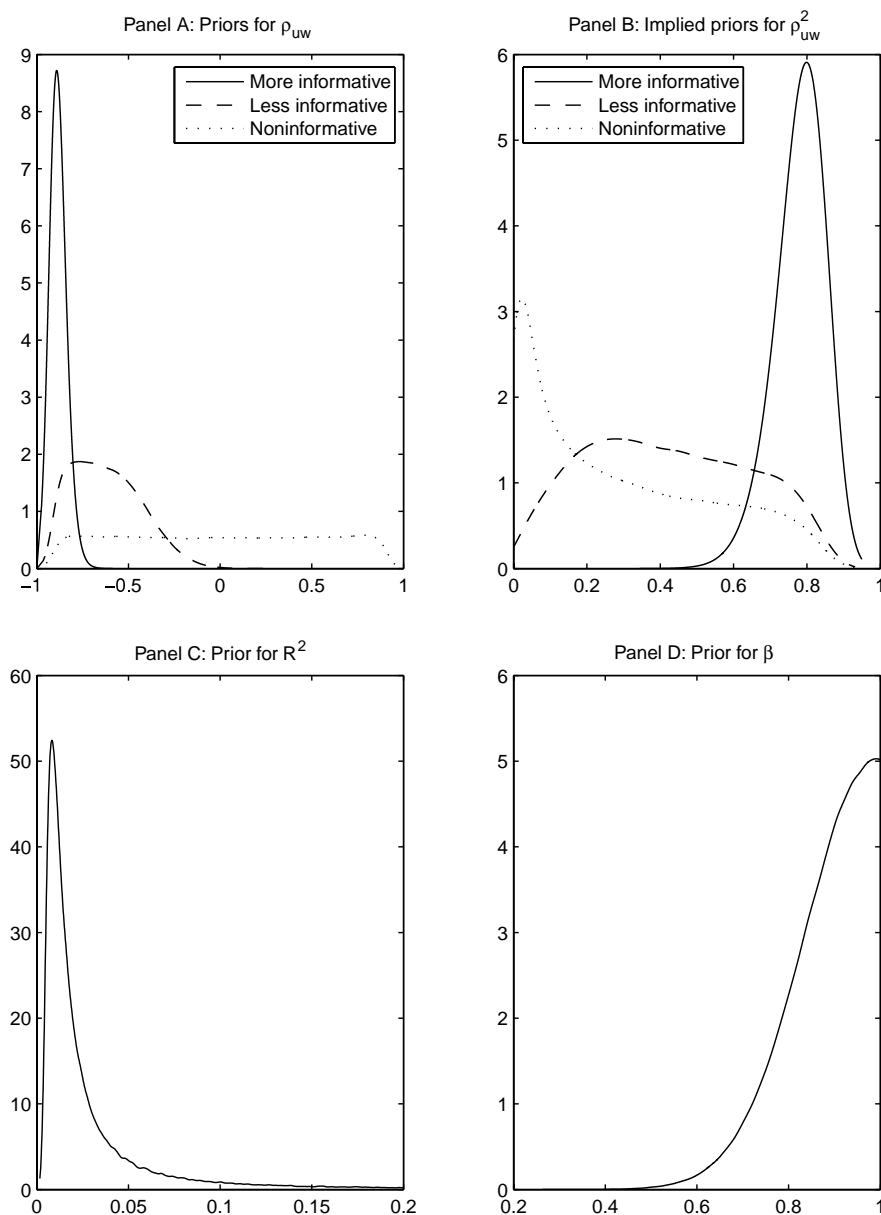
Figure 5. Prior distributions. Panel A plots three prior distributions for the correlation between expected and unexpected returns, ρ_{uw}. The noninformative prior (dotted line) is flat between -0.9 and 0.9, with tails fading away as ρ_{uw} approaches ± 1. The less informative prior (dashed line) has 99.9% of its mass below zero ($\rho_{uw} < 0$). The more informative prior (solid line) has 99.9% of its mass below -0.71, so that $\rho_{uw}^2 > 0.5$ (i.e., unexpected changes in the discount rate explain over half of the variance of unexpected market returns). Panel B plots the corresponding implied priors on ρ_{uw}^2. Panel C plots the prior on the predictive R^2 from the regression of returns r_{t+1} on expected returns μ_t. Panel D plots the prior on the slope coefficient β in the AR(1) process for μ_t. All parameters correspond to quarterly data.

The three priors on ρ_{uw} are chosen to represent a wide range of possible prior beliefs. The noninformative prior is obviously weak, whereas the more informative prior can probably be characterized as strong. For example, the latter prior is slightly stronger than the posteriors of ρ_{uw} based on various estimates of Campbell (1991) from a sample period that predates ours: Campbell's evidence from 1927 to 1951 implies estimates of ρ_{uw}^2 ranging from 0.44 to 0.76 (i.e., ρ_{uw} ranging from -0.67 to -0.87). The more informative prior does not appear to be extreme, though, as it is consistent with most empirical estimates. For example, the evidence of Campbell (1991) implies estimates of ρ_{uw}^2 ranging from 0.50 to 0.74 across three different specifications in his full sample period 1927 to 1988. The estimates of ρ_{uw}^2 implied by the post-war evidence range from 0.84 to 0.88 in 1952 to 1988, and the estimates of Campbell and Ammer (1993) in their Table III range from 0.86 to 0.91. All of these estimates are in line with the more informative prior.

Putting a prior on ρ_{uw} presents a technical challenge. We do not impose the standard inverted Wishart prior on the covariance matrix Σ because such a prior would be informative about all elements of Σ, not only about ρ_{uw}, and we do not wish to be informative about the elements that involve v_t. Instead, we build on Stambaugh (1997) and form the prior on Σ as the posterior from a hypothetical sample that contains more information about u_t and w_t than about v_t. In addition, we develop a hyperparameter approach that allows us to change the prior on ρ_{uw} without changing the priors on any other parameters. Further discussion of the priors appears in the Appendix.

In addition to putting a prior on ρ_{uw}, we also impose a prior belief that the conditional expected return μ_t is stable and persistent. To capture the belief that μ_t is stable, we impose a prior that the predictive R^2 from the regression of r_{t+1} on μ_t is not very large, which is equivalent to the belief that the total variance of μ_t is not very large. The prior on the R^2, which is plotted in Panel C of Figure 5, has a mode close to 1%, most of its mass is below 5%, and there is very little prior mass above 10%. To capture the belief that μ_t is persistent, we impose a prior that β, the slope of the AR(1) process for μ_t, is smaller than one but not by much.[17] The prior on β, which is plotted in Panel D of Figure 5, has most of its mass above 0.7 and there is virtually no prior mass below 0.5. We do not impose a prior belief that $\mu_t > 0$. Although such a belief is reasonable under a fully rational view, we do not wish to preclude the possibility that some of the variation in μ_t is driven by investor sentiment. The prior distributions on all other parameters (E_x, A, E_r, and most elements of Σ) are noninformative. Separately, we also consider a "diffuse" prior, which is completely noninformative about all model parameters, including ρ_{uw}, β, and R^2.

D. Predictive System versus Predictive Regression

In contrast to a predictive regression, the predictive system allows us to conduct finite-sample inferences that explicitly incorporate predictor imperfection.

---

[17] Ferson, Sarkissian, and Simin (2003, footnote 2) discuss several reasons to believe expected return is persistent.

The predictive system also produces more precise inferences about expected returns. To demonstrate this, we compare the explanatory powers of the system and the regression for a broad range of parameter values. Specifically, we compare the R^2 in the regression of r_{t+1} on x_t for the predictive regression with the R^2 in the regression of r_{t+1} on $\mathrm{E}(r_{t+1} \mid D_t) \equiv \mathrm{E}(\mu_t \mid D_t)$ for the predictive system. The ratio of these R^2 values when r_{t+1} is the dependent variable is the same as when μ_t is the dependent variable,

$$\frac{R^2(r_{t+1} \text{ on } x_t)}{R^2(r_{t+1} \text{ on } \mathrm{E}(\mu_t \mid D_t))} = \frac{R^2(\mu_t \text{ on } x_t)}{R^2(\mu_t \text{ on } \mathrm{E}(\mu_t \mid D_t))}, \tag{29}$$

since each of the R^2 values in the latter ratio is equal to its corresponding value in the first ratio multiplied by $\mathrm{Var}(r_{t+1})/\mathrm{Var}(\mu_t)$. The parameters in equations (3) to (6) can be used to obtain the covariance matrix of μ_t and x_t and thereby the R^2 in the regression of μ_t on x_t,

$$R^2(reg) = \frac{\mathrm{Var}[\mathrm{E}(\mu_t \mid x_t)]}{\mathrm{Var}(\mu_t)}. \tag{30}$$

As shown in the Appendix, we can solve analytically for the steady-state value of $\mathrm{Var}(\mu_t \mid D_t)$, which allows us to compute the R^2 in the regression of μ_t on $\mathrm{E}(\mu_t \mid D_t)$ as

$$R^2(sys) = \frac{\mathrm{Var}[\mathrm{E}(\mu_t \mid D_t)]}{\mathrm{Var}(\mu_t)} = 1 - \frac{\mathrm{Var}(\mu_t \mid D_t)}{\mathrm{Var}(\mu_t)}. \tag{31}$$

The ratio in equation (29) is computed as $R^2(reg)/R^2(sys)$. Note that this R^2 ratio cannot exceed one because $x_t \in D_t$. In other words, the estimates of μ_t from the predictive system are at least as precise as the estimates from the predictive regression, simply because the system uses more information. The smaller the R^2 ratio, the larger the advantage of using the predictive system.

We use the R^2 ratios to quantify the explanatory advantage of the predictive system, using the same sample as in Section III.A. Panel A of Table II shows the posterior means and *SD* of the R^2 ratios for four different priors and four different sets of predictors. First, observe that the posterior means of the R^2 ratios are all comfortably lower than one, ranging from 0.08 to 0.86 across the 16 cases, and from 0.46 to 0.70 when all three predictors are used jointly. Second, the R^2 ratios are sensitive to the prior on ρ_{uw}. For example, with the bond yield as the single predictor, the R^2 ratio is estimated to be 0.73 under the diffuse prior. When we impose the prior belief that ρ_{uw} is negative, the R^2 ratio declines to 0.34 under the less informative prior and then further to 0.08 under the more informative prior. In other words, under the prior that more than half of the market variance is due to discount rate news, the expected return estimates from the predictive system are about 12.5 times more precise than those from the predictive regression. For the dividend yield, we observe the opposite pattern—the R^2 ratio increases from 0.28 to 0.59 to 0.81 for the same priors. The opposite patterns result from the opposite effects that the prior on ρ_{uw} has on the adequacy of x_t as a predictor in the two cases, as we will see later.

Table II
Explanatory Power of the Predictive Regression Relative to the Predictive System

Panel A shows the posterior means and standard deviations (the latter in parentheses) of the ratios of two R^2s, $R^2(reg)/R^2(sys)$. A ratio smaller than one indicates that the predictive system estimates μ_t more precisely than the predictive regression does. The smaller the ratio, the larger the advantage of using the predictive system. $R^2(reg)$, computed as the R^2 from the regression of μ_t on the given predictors, summarizes the usefulness of the predictive regression in estimating μ_t. $R^2(sys)$, computed as $1 - \mathrm{Var}(\mu_t \mid D_t)/\mathrm{Var}(\mu_t)$, where D_t contains all historical market returns and predictor realizations, summarizes the usefulness of the predictive system in estimating μ_t. Panel B shows the posterior means and *SD* of $1 - MSE(sys)/MSE(reg)$. $MSE(reg)$ is the mean squared error from the predictive regression of r_{t+1} on the given predictors. $MSE(sys)$ is the mean squared error from the predictive system. Positive values of one minus the MSE ratio indicate that the predictive system forecasts returns more precisely than the predictive regression does. The results are reported for four different prior distributions on ρ_{uw}, the correlation between expected and unexpected returns. Four sets of predictors are considered: dividend yield, bond yield, CAY, and all three predictors combined. The sample period is 1952Q1 to 2003Q4.

	Predictors			
	Dividend Yield	Bond Yield	CAY	All 3 Predictors
Panel A: The R^2 Ratios, $R^2(reg)/R^2(sys)$				
Diffuse	0.28	0.73	0.86	0.59
Prior	(0.17)	(0.23)	(0.16)	(0.30)
Noninformative	0.50	0.44	0.61	0.46
Prior on ρ_{uw}	(0.27)	(0.25)	(0.27)	(0.22)
Less Informative	0.59	0.34	0.73	0.50
Prior on ρ_{uw}	(0.22)	(0.20)	(0.23)	(0.22)
More Informative	0.81	0.08	0.64	0.70
Prior on ρ_{uw}	(0.19)	(0.08)	(0.22)	(0.19)
Panel B: The Mean Squared Error Ratios, $1 - MSE(sys)/MSE(reg)$				
Diffuse	0.04	0.06	0.03	0.17
Prior	(0.07)	(0.15)	(0.09)	(0.22)
Noninformative	0.02	0.02	0.02	0.03
Prior on ρ_{uw}	(0.04)	(0.05)	(0.04)	(0.04)
Less Informative	0.02	0.02	0.01	0.04
Prior on ρ_{uw}	(0.04)	(0.06)	(0.04)	(0.05)
More Informative	0.02	0.03	0.02	0.02
Prior on ρ_{uw}	(0.07)	(0.07)	(0.06)	(0.07)

Panel B of Table II shows the posterior means and *SD* of one minus the ratio of the mean squared errors from the predictive system and the predictive regression. The mean squared error is defined as $MSE = \mathrm{E}\{(r_{t+1} - f_t)^2\}$, where f_t is a return forecast. All posterior means in Panel B are positive, ranging from 0.01 to 0.17, confirming that the predictive system forecasts returns more precisely than the predictive regression does.

Another way of comparing the predictive system with the predictive regression is to compare their estimates of the slope coefficient b from the predictive
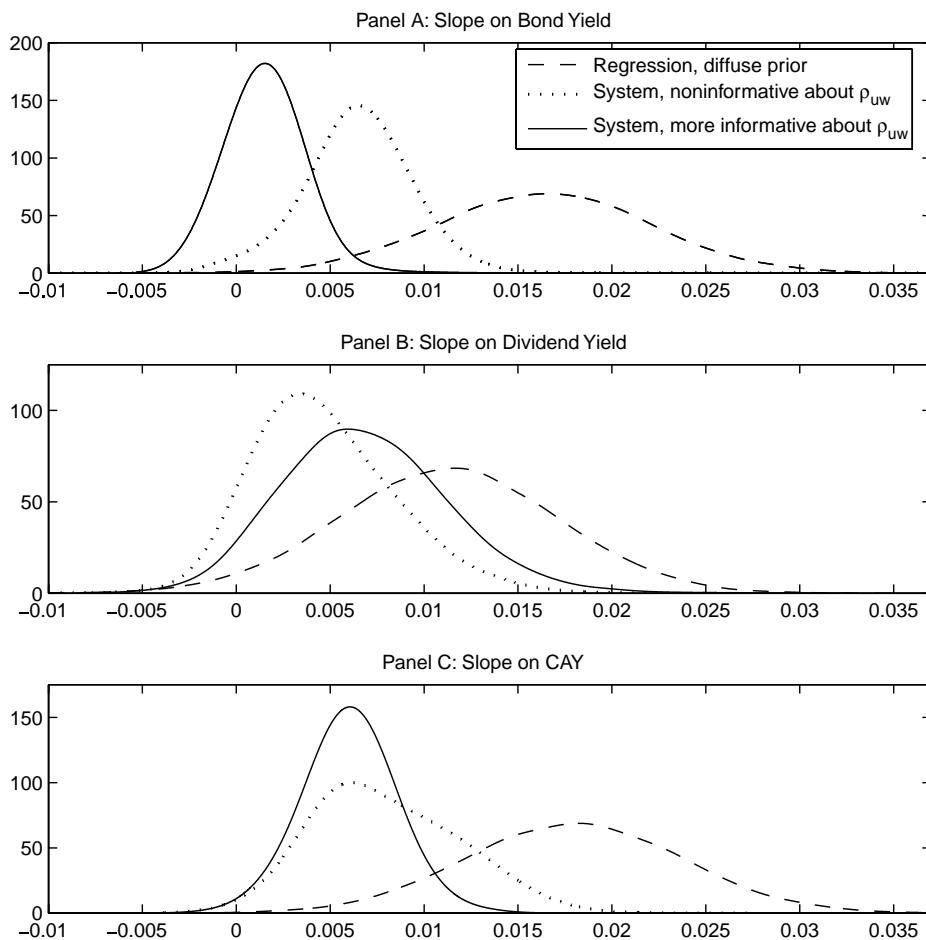
Figure 6. Posterior distributions of slope coefficients from predictive regressions. We estimate both the predictive system and the predictive regression with three predictors: the bond yield, dividend yield, and CAY. The dashed line plots the posteriors from the standard predictive regression of r_{t+1} on x_t under the diffuse prior. The dotted line plots the implied posteriors constructed from the results of the predictive system under the "noninformative" prior on ρ_{uw}. The solid line plots the implied posteriors from the predictive system under the "more informative" prior on ρ_{uw} ($\rho_{uw} < -0.71$). To facilitate comparisons across panels, all predictors are scaled to have unit variance. The sample period is 1952Q1 to 2003Q4.

regression. Figure 6 plots the posterior distributions of b computed under three scenarios. The dashed line is the posterior of b computed from the predictive regression under no prior information. This posterior has a Student t distribution whose mean is equal to the maximum likelihood estimate (MLE) of b (Zellner (1971), pp. 65–67). The dashed line thus represents "conventional inference" on predictability. The other two lines in Figure 6 plot the implied

posteriors of b computed from the predictive system.[18] The dotted line corresponds to the prior that is noninformative about ρ_{uw} but informative about β and R^2. In all three panels of Figure 6, the dotted line is substantially different from the dashed line, which means that imposing the prior that μ_t is stable and persistent significantly affects the inference about predictability. In addition, the dotted line is shifted toward zero compared to the dashed line, which means that the prior belief that μ_t is stable and persistent weakens the evidence of predictability. Finally, the solid line corresponds to the prior that is informative not only about β and R^2 but also about ρ_{uw}. The prior on ρ_{uw} clearly affects the inference about predictability. Consider Panel A, in which the single predictor is the bond yield. Whereas the traditional inference (dashed line) would conclude with almost 100% certainty that the bond yield is a useful predictor ($b > 0$), the system-based inference with the more informative prior on ρ_{uw} (solid line) concludes no such thing because almost half of the posterior mass of b is below zero. This prior also slightly weakens the predictive power of CAY but it strengthens the predictive power of the dividend yield.

E. How Imperfect Are Predictors?

The predictive system also allows us to learn about the correlation between the expected return μ_t and the predictors. Since μ_t is not observed, the manner in which one learns about such correlations merits some discussion. Consider, for simplicity, the case of a single predictor x_t whose autocorrelation A is equal to β. The unconditional correlation between the expected return and the predictor, $\rho_{x\mu}$, is then equal to ρ_{vw}, the conditional correlation.[19] By virtue of the fact that the correlation matrix for (u_t, v_t, w_t) must be nonnegative definite, it is readily verified that

$$\rho_{vw} = \rho_{uv}\rho_{uw} + \rho_\Delta, \quad \text{where} \quad \rho_\Delta^2 \le \left(1 - \rho_{uv}^2\right)\left(1 - \rho_{uw}^2\right). \tag{32}$$

In other words, even though correlations are not transitive (two correlations do not imply the third), they become nearly transitive when at least one of them approaches ± 1.

We specify noninformative priors for ρ_{vw} and ρ_{uv}. Doing otherwise would involve priors about each predictor's usefulness—directly through ρ_{vw} but indirectly through ρ_{uv} as well, given informative priors about ρ_{uw}. While such an approach could be reasonable, especially in a forecasting setting, we wish to illustrate how our framework can deliver inferences about each predictor's usefulness without making prior judgments about that property. Moreover, we

---

[18] Although b does not appear explicitly in the predictive system, its value can be computed from the system's parameters (see footnote 6), so its posterior draws can be constructed from the draws of the system's parameters.

[19] More generally,

$$\rho_{x\mu} = \rho_{vw} \left[\frac{(1 - \beta^2)(1 - A^2)}{(1 - \beta A)^2} \right]^{\frac{1}{2}},$$

so that $\rho_{x\mu}^2 \le \rho_{vw}^2$.

prefer not to add complexity to this initial exploration of predictive systems. The data are quite informative about ρ_{uv} anyway, in that with only modest return predictability, the value of ρ_{uv} is close to that of ρ_{rv}, which can be estimated from the series of r_t and x_t. (When the predictive R^2 is low, $\rho_{ru}^2 = (1 - R^2)$ is close to one, and equation (32) then implies that ρ_{uv} is well approximated by ρ_{rv}.) Information about ρ_{uw} enters largely through the prior. When the prior is concentrated on large negative values, then the likely values of ρ_{Δ} in (32) are small, so the prior information about ρ_{uw} and the sample information about ρ_{uv} get combined to provide information about ρ_{vw}. Alternatively, if the data indicate that ρ_{uv} is close to ± 1 (e.g., in Table I, $\rho_{uv} \approx -0.9$ when the predictor is the dividend yield), then again ρ_{Δ} is small, so that ρ_{uv} and ρ_{uw} are again jointly informative about ρ_{vw}.

Figures 7 and 8 analyze the degree to which two of our predictors can capture the unobservable expected return μ_t. We report results for two predictive systems, in which the predictors are the dividend yield alone (Figure 7) and the bond yield alone (Figure 8). Panel A of each figure plots the posterior distribution of the R^2 from the regression of μ_t on x_t. This R^2 is assumed to be one in a predictive regression, but its posterior in the predictive system has very little mass at values close to one. In both figures, the R^2s larger than 0.8 receive very little posterior probability and the values larger than 0.9 are deemed almost impossible, regardless of the prior. This evidence suggests that neither of the two predictors is perfectly correlated with μ_t.[20]

The R^2 depends on the prior for ρ_{uw} in an interesting way. In Panel A of Figure 7, becoming increasingly informative about ρ_{uw} shifts the posterior of the R^2 to the right, with the mode shifting from about 0.3 under the noninformative prior to about 0.6 under the more informative prior. This makes sense: Since the dividend yield exhibits a highly negative contemporaneous correlation with stock returns, imposing a prior that μ_t also possesses such negative correlation makes the dividend yield more closely related to μ_t. Exactly the opposite happens in Panel A of Figure 8, where becoming increasingly informative about ρ_{uw} shifts the posterior of the R^2 to the left so that its mode is close to zero under the more informative prior. This makes sense as well because the bond yield is positively correlated with stock returns (Table I).

Panel B of both figures plots the posterior of the predictive R^2 from the regression of r_{t+1} on μ_t. Putting a more informative prior on ρ_{uw} increases the R^2 in both figures, but these effects are relatively small. Since we put a fairly informative prior on the predictive R^2 (see Panel C of Figure 5), the posterior is not dramatically different from the prior in either figure.

Panels C and D of both figures plot the posteriors of the correlations between each predictor and μ_t, both conditional (ρ_{vw}) and unconditional ($\rho_{x\mu}$). These correlations are all well below 1 and they are quite sensitive to the prior on ρ_{uw}. As we become increasingly informative about ρ_{uw}, we perceive the dividend

---

[20] Note that even if x_t were perfectly correlated with μ_t in population, the posterior of their correlation would have nontrivial mass below one in any finite sample. Since we always observe finite samples, we always perceive imperfect correlation between x_t and μ_t. Also note that in the NBER version of this article, we report results analogous to those in Figures 7 and 8 for a predictive system that uses three predictors: the dividend yield, bond yield, and CAY.
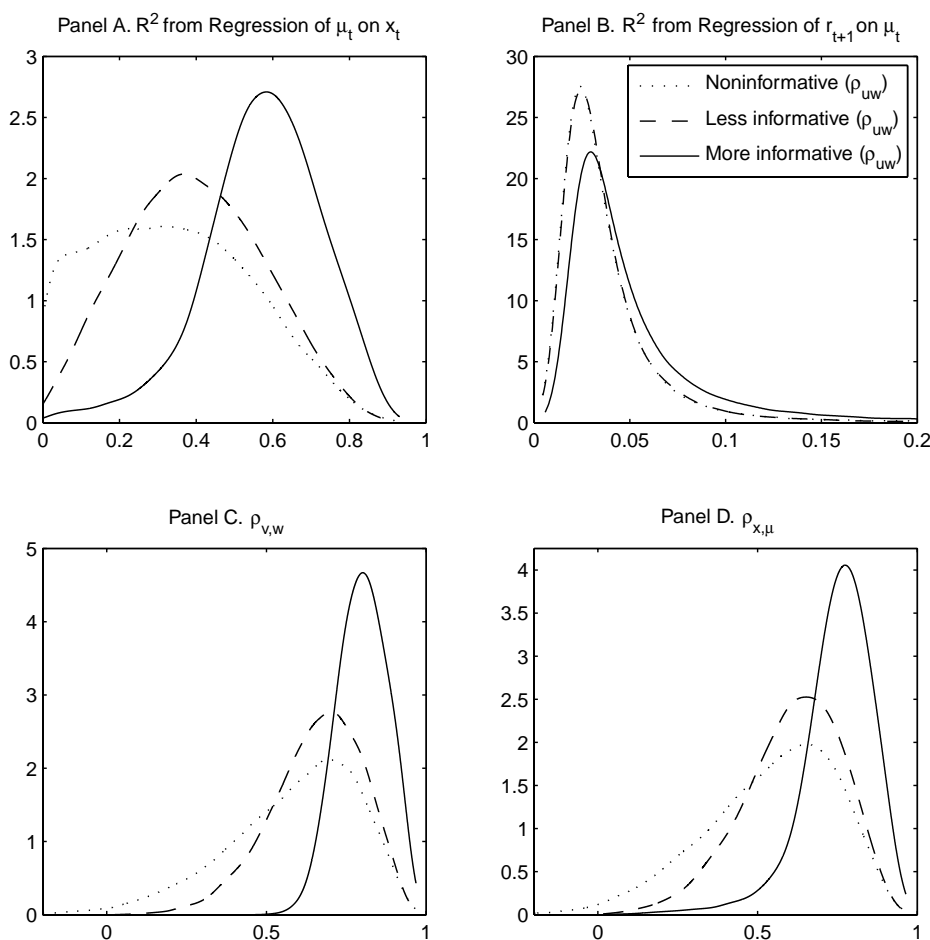
Panel A. R² from Regression of μ_t on x_t

Panel B. R² from Regression of r_{t+1} on μ_t

Panel C. $\rho_{v,w}$

Panel D. $\rho_{x,\mu}$

Figure 7. Posterior distributions for the relations between the dividend yield and expected return. Panel A plots the posterior of the fraction of variation in the expected return μ_t that can be explained by the predictor x_t, which is the dividend yield. Panel B plots the posterior of the predictive R^2. Panel C plots the posterior of the *conditional* correlation ρ_{vw} between the dividend yield and μ_t. Panel D plots the posterior of the *unconditional* correlation $\rho_{x\mu}$ between the dividend yield and μ_t. The three lines in each panel represent three different prior distributions. The solid line represents the "more informative" prior on $\rho_{uw}(\rho_{uw} < -0.71)$, the dashed line is the "less informative" prior on $\rho_{uw}(\rho_{uw} < 0)$, and the dotted line is the "noninformative" prior on ρ_{uw}. The sample period is 1952Q1 to 2003Q4.

yield to be more highly correlated with μ_t but the bond yield to be less highly correlated with μ_t. The effect for the bond yield is dramatic, judging by the posterior modes in Panel C of Figure 8. Under the noninformative prior, the bond yield has 90% conditional correlation with μ_t, but under the more informative prior, this correlation drops to 20%.

Overall, Figures 7 and 8 show that our predictors are imperfectly correlated with μ_t and that the inference about this correlation is substantially affected by
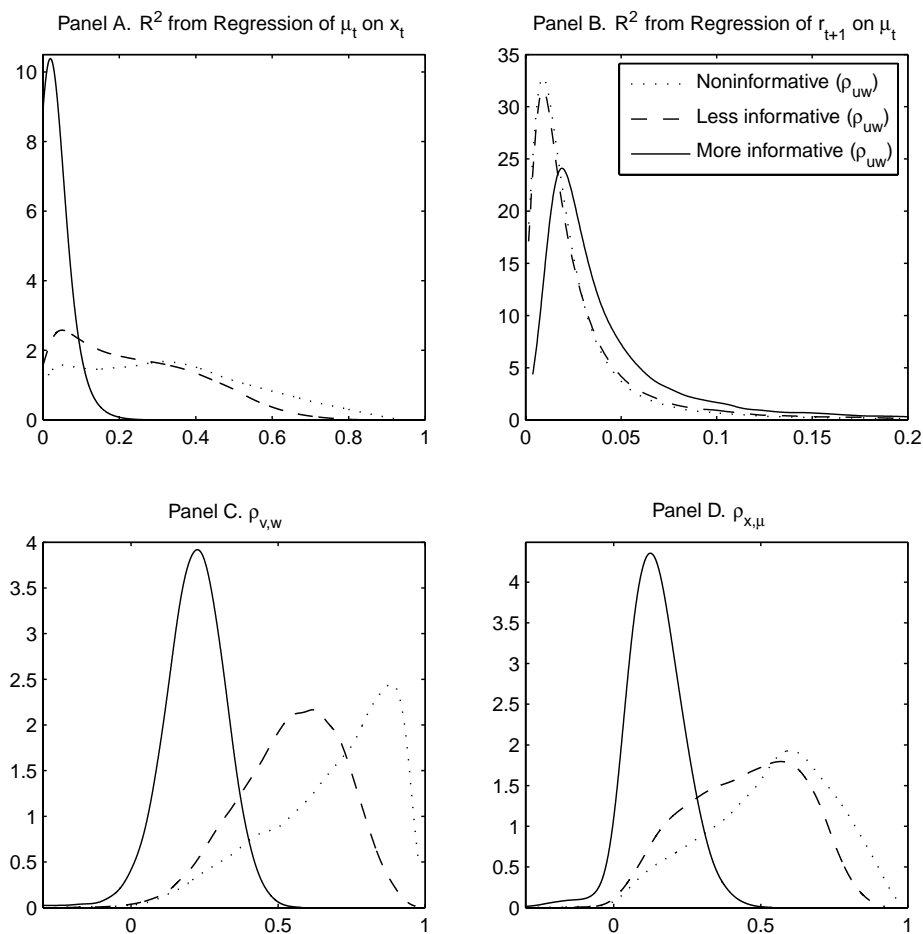
Figure 8. Posterior distributions for the relations between the bond yield and expected return. Panel A plots the posterior of the fraction of variation in the expected return μ_t that can be explained by the predictor x_t, which is the bond yield. Panel B plots the posterior of the predictive R^2. Panel C plots the posterior of the *conditional* correlation ρ_{vw} between the bond yield and μ_t. Panel D plots the posterior of the *unconditional* correlation $\rho_{x\mu}$ between the bond yield and μ_t. The three lines in each panel represent three different prior distributions. The solid line represents the "more informative" prior on $\rho_{uw}(\rho_{uw} < -0.71)$, the dashed line is the "less informative" prior on $\rho_{uw}(\rho_{uw} < 0)$, and the dotted line is the "noninformative" prior on ρ_{uw}. The sample period is 1952Q1 to 2003Q4.

the prior beliefs about ρ_{uw}. Prior beliefs that $\rho_{uw} < 0$ strengthen the predictive appeal of the dividend yield but they weaken the predictive appeal of the bond yield.

F. Estimates of Expected Return

Figure 9 plots the time series of expected returns estimated by three different approaches. The dashed line plots the fitted values from the predictive
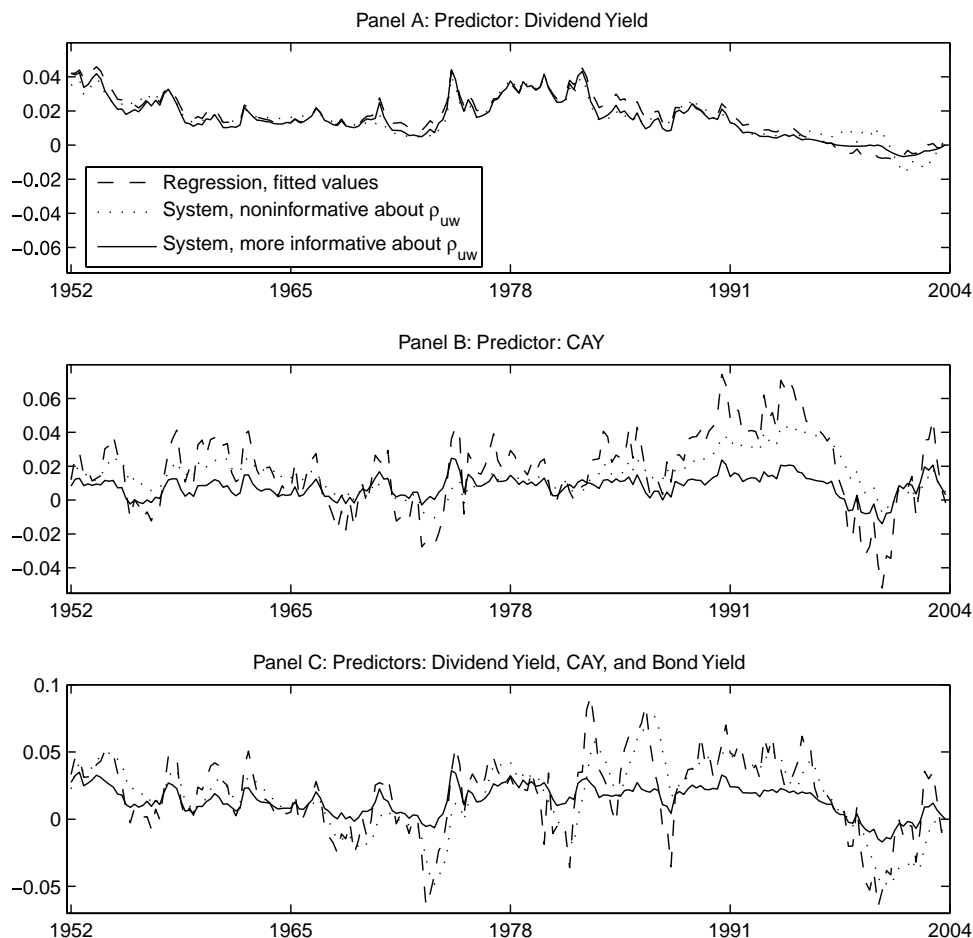
Figure 9. The equity premium: Regression versus system with prior information. This figure plots the time series of the quarterly equity premium estimated in three different environments. The dashed line plots the OLS fitted values from the predictive regression of r_{t+1} on the given predictor(s). The dotted line plots the posterior means of $E(r_{t+1} | D_t)$ from the predictive system under the "noninformative" prior on ρ_{uw}. The solid line plots the posterior means of $E(r_{t+1} | D_t)$ from the predictive system under the "more informative" prior on $\rho_{uw} (\rho_{uw} < -0.71)$. In Panel A, the estimation uses one predictor, dividend yield. In Panel B, the single predictor is CAY. In Panel C, three predictors are used: dividend yield, CAY, and bond yield. The sample period is 1952Q1 to 2003Q4.

regression. These traditional expected return estimates seem too volatile to be plausible, as we also observed in Figure 4. For example, in Panel C, which includes all three predictors, expected returns range from −6% to 9% per quarter. Not surprisingly, imposing the prior that μ_t is stable and persistent (dotted line) produces smoother expected return estimates. Adding the more informative prior on ρ_{uw} (solid line) further smoothes the expected return estimates: in Panel C, they range from −1.5% to 3.5% per quarter. The informative priors

have substantial effects on expected returns not only in Panel C but also in Panel B in which CAY is the single predictor: while the regression-fitted values range from -5.5% to 7.5% per quarter, the solid line ranges from -1.5% to 2.5%. Only in Panel A, in which the dividend yield is the single predictor, is the effect of the prior relatively mild. The reason is that the regression-fitted values in Panel A are already fairly smooth and negatively correlated with stock returns.

While eyeballing the expected return estimates seems informative, we also compute measures summarizing their differences. Table III compares five different series of expected return estimates. The first is the series of fitted values from the predictive regression, and the others are produced by four different approaches to estimating the predictive system. One of the latter approaches estimates the predictive system by MLE, while the other three impose the prior that μ_t is stable and persistent but differ in their prior on ρ_{uw}. We compare the five series of expected return estimates in three different ways: pairwise correlations, mean absolute differences, and average utility losses. The utility losses are computed for a mean-variance investor allocating between the market and the T-bill who knows the variance of market returns but must estimate the market's expected return. The investor's risk aversion, 2.54, is such that the optimal portfolio is fully invested in the market, on average. We compute the investor's certainty equivalent loss resulting from holding a portfolio that is optimal under a different approach for estimating expected returns. For example, the 0.11% per quarter average utility loss in the first row of Panel A is suffered by an investor who wants to estimate expected return in the predictive system by MLE but is forced to use the fitted values from the predictive regression. Finally, the three panels consider three different sets of predictors: the dividend yield, CAY, and the two predictors combined with the bond yield.

Panel A of Table III shows that when the dividend yield is the single predictor, the expected return estimates are fairly similar across the five estimation approaches, confirming the evidence from Panel A of Figure 9. No average utility loss exceeds 0.20% per quarter, no mean absolute difference is larger than 0.65% per quarter, and all correlations exceed 81%.

The differences across the five approaches are substantially larger in Panel B where we use CAY to predict returns. For example, compare the system-based estimates obtained by MLE versus the more informative prior. The mean absolute difference in expected returns is 1.65% per quarter, and the average certainty equivalent loss from using one estimate in place of the other is 1.38% per quarter. Both quantities are highly economically significant. In Panel C, where we use all three predictors, the differences across the five approaches are also large and similar in magnitude.

In all three panels, the smallest differences are obtained for the noninformative versus the less informative prior on ρ_{uw}. No average utility loss exceeds 0.06% per quarter, no mean absolute difference is larger than 0.37% per quarter, and all correlations exceed 95.4%. However, moving from the less informative to the more informative prior on ρ_{uw} can produce sizeable differences in expected returns. For example, the mean absolute difference in Panel C is 1.46% per quarter and the average utility loss is 0.84% per quarter.

Table III

Comparing Estimates of Expected Return

This table compares the time series of the posterior means of $E(r_{t+1} | D_t)$ obtained in five different environments:

(1) Predictive regression: OLS fitted values.
(2) Predictive system: Maximum likelihood estimates.
(3) Predictive system: Noninformative prior about ρ_{uw}.
(4) Predictive system: Less informative prior about ρ_{uw}.
(5) Predictive system: More informative prior about ρ_{uw}.

The priors in (3)–(5) are informative about the persistence and volatility of μ_t. The correlations between the quarterly series of the posterior means of $E(r_{t+1} | D_t)$ are reported in italics below the main diagonal of each left-panel 5×5 matrix. Above the main diagonal of the same matrix are the mean absolute differences between the posterior means of $E(r_{t+1} | D_t)$ in percent per quarter. Each right-panel 5×5 matrix reports the average utility losses, in percent per quarter, of a mean-variance investor who is forced to hold a suboptimal portfolio of the stock market and a risk-free T-bill: a portfolio that is optimal under the beliefs in the given row when the true beliefs are in the given column. (For example, the (2,5) cell of the 5×5 matrix reports the certainty equivalent loss of an investor who has the more informative prior but is forced to hold the portfolio that is optimal under the maximum likelihood estimates.) The risk aversion is chosen such that there is no borrowing or lending given the sample mean and variance of market returns. The sample period is 1952Q1 to 2003Q4.

	Correlation (%)/Mean Abs Diff (%)					Average Utility Loss (%)				
	(1)	(2)	(3)	(4)	(5)	(1)	(2)	(3)	(4)	(5)
Panel A: Predictor: Dividend Yield										
(1)		0.57	0.44	0.41	0.29	0	0.11	0.09	0.07	0.03
(2)	*97.49*		0.65	0.62	0.49	0.11	0	0.20	0.18	0.10
(3)	*90.47*	*81.11*		0.06	0.27	0.09	0.19	0	0.00	0.03
(4)	*92.42*	*83.49*	*99.78*		0.22	0.07	0.17	0.00	0	0.02
(5)	*97.67*	*91.32*	*95.81*	*97.41*		0.03	0.10	0.03	0.02	0
Panel B: Predictor: CAY										
(1)		0.66	1.13	1.22	1.60	0	0.19	0.57	0.65	1.13
(2)	*94.82*		1.36	1.39	1.65	0.19	0	0.88	0.95	1.38
(3)	*86.28*	*82.02*		0.37	0.96	0.57	0.85	0	0.06	0.39
(4)	*96.88*	*93.26*	*95.43*		0.63	0.64	0.92	0.06	0	0.16
(5)	*89.71*	*92.03*	*59.38*	*80.11*		1.09	1.32	0.38	0.16	0
Panel C: Predictors: Dividend Yield, CAY, Bond Yield										
(1)		0.92	1.33	1.27	1.60	0	0.40	0.82	0.74	1.19
(2)	*91.06*		1.27	1.22	1.62	0.42	0	0.76	0.68	1.16
(3)	*80.38*	*82.36*		0.14	1.51	0.80	0.72	0	0.01	0.94
(4)	*82.30*	*84.11*	*99.79*		1.46	0.72	0.64	0.01	0	0.84
(5)	*83.42*	*89.93*	*80.75*	*84.00*		1.19	1.13	0.96	0.87	0

To sum up, when we use the dividend yield as the single predictor, the system-based expected return estimates are close to the regression-based estimates. In all other cases, the system and the regression generate substantially different expected returns, and the system-based estimates are significantly affected by the prior on ρ_{uw}.

G. Variance Decomposition of Expected Return

In the predictive regression approach, expected return μ_t is modeled as an exact linear function of the predictors in x_t. In a predictive system, however, the data provide additional information about μ_t because the lagged values of unexpected returns and predictor innovations also enter the expected return estimates (see Section II.B). In this section, we decompose the variance of μ_t to assess the relative importance of the various sources of information in a predictive system.

We can rewrite the AR(1) process for x_t as an MA(∞) process, as we did for μ_t in equation (14): $x_t = E_x + \sum_{i=0}^{\infty} A^i v_{t-i}$. Then we project w_t linearly on u_t and v_t:

$$w_t = [\sigma_{wu} \quad \sigma_{wv}] \begin{bmatrix} \sigma_u^2 & \sigma_{uv} \\ \sigma_{vu} & \Sigma_{vv} \end{bmatrix}^{-1} \begin{bmatrix} u_t \\ v_t \end{bmatrix} + \chi_t = \psi_u u_t + \psi_v v_t + \chi_t. \tag{33}$$

Substituting for w_t from equation (33) into equation (14), we obtain

$$\mu_t = (E_r - \psi_v E_x) + \psi_v x_t + \psi_u \sum_{i=0}^{\infty} \beta^i u_{t-i} + \psi_v \sum_{i=0}^{\infty} (\beta^i I_K - A^i) v_{t-i} + \sum_{i=0}^{\infty} \beta^i \chi_{t-i}, \tag{34}$$

where K is the number of predictors and I_K is a $K \times K$ identity matrix. Equation (34) shows how the lagged values of unexpected returns u_{t-i} and predictor innovations v_{t-i} affect μ_t in the presence of the current predictor values in x_t. Based on this equation, we can decompose the variance of μ_t into the components due to x_t, $\{u_s\}_{s \leq t}$, and $\{v_s\}_{s \leq t}$. See the Appendix for details.

Table IV reports the posterior means and *SD* of the R^2s from the regressions of μ_t on x_t (column (1)), μ_t on x_t and $\{u_s\}_{s \leq t}$ (column (2)), and μ_t on x_t and $\{u_s, v_s\}_{s \leq t}$ (column (3)). For x_t, we consider four sets of predictors: the dividend yield, bond yield, CAY, and the combination of all three predictors. For each set of predictors, we estimate the predictive system under three different priors. All three priors assume that μ_t is stable and persistent but they differ in their degree of informativeness about ρ_{uw}.

First, note that x_t never accounts for more than 63% of the variance of μ_t and that it can account for as little as 3% of this variance. In contrast, x_t combined with $\{u_s, v_s\}_{s \leq t}$ can account for as much as 95% of the variance of μ_t, and those components account for more than 80% of the variance of μ_t in 10 of the 12 cases in Table IV. The most striking effect obtains for the bond yield, for which adding $\{u_s, v_s\}_{s \leq t}$ to x_t increases the R^2 from 0.03 to 0.95. It seems clear that a predictive regression, which uses only x_t to predict returns, does not use the data as effectively as a predictive system, which also uses $\{u_s, v_s\}_{s \leq t}$ in addition to x_t.

The R^2s in Table IV are substantially affected by the prior on ρ_{uw}. For example, consider the first columns of Panels A and B. Under the noninformative prior on ρ_{uw}, both the dividend yield and the bond yield explain about a third of the variance of μ_t. As we become more informative about ρ_{uw}, this fraction increases from 0.34 to 0.40 to 0.57 for the dividend yield, but it decreases from

<div align="center">

Table IV

Variance Decomposition of Expected Return
</div>

This table reports the posterior means and *SD*s (the latter in parentheses) of the R^2s from the regressions of the market's expected excess return μ_t on its selected components. The first column of each panel, labeled x_t, shows the fraction of variance of μ_t that can be explained by the set of predictors listed in the panel heading. Four sets of predictors are considered: the dividend yield, bond yield, CAY, and the combination of all three of these predictors. The second column of each panel, labeled $x_t, \{u_s\}_{s\leq t}$, shows the fraction of variance of μ_t that can be explained jointly by the predictors and by the innovations to stock market returns $u_t, u_{t-1}, u_{t-2}, \ldots$. The third column, labeled $x_t, \{u_s, v_s\}_{s\leq t}$, shows the fraction of variance of μ_t that can be explained jointly by the predictors, by the innovations to stock market returns $u_t, u_{t-1}, u_{t-2}, \ldots$, and by the innovations to the predictors $v_t, v_{t-1}, v_{t-2}, \ldots$. For each set of predictors, the predictive system is estimated under three different priors, which are described in the row labels. The sample period is 1952Q1 to 2003Q4.

	Components of Expected Return			Components of Expected Return		
	x_t	$x_t, \{u_s\}_{s\leq t}$	$x_t, \{u_s, v_s\}_{s\leq t}$	x_t	$x_t, \{u_s\}_{s\leq t}$	$x_t, \{u_s, v_s\}_{s\leq t}$
	Panel A: Dividend Yield			Panel B: Bond Yield		
Noninformative Prior on ρ_{uw}	0.34 (0.20)	0.43 (0.21)	0.48 (0.21)	0.33 (0.21)	0.64 (0.21)	0.83 (0.13)
Less informative Prior on ρ_{uw}	0.40 (0.18)	0.49 (0.18)	0.53 (0.18)	0.24 (0.17)	0.73 (0.16)	0.86 (0.11)
More informative Prior on ρ_{uw}	0.57 (0.15)	0.80 (0.06)	0.81 (0.06)	0.03 (0.03)	0.86 (0.05)	0.95 (0.04)
	Panel C: CAY			Panel D: All Three Predictors		
Noninformative Prior on ρ_{uw}	0.50 (0.22)	0.59 (0.22)	0.81 (0.12)	0.42 (0.20)	0.49 (0.22)	0.90 (0.07)
Less informative Prior on ρ_{uw}	0.60 (0.20)	0.70 (0.17)	0.83 (0.10)	0.46 (0.20)	0.55 (0.22)	0.90 (0.07)
More informative Prior on ρ_{uw}	0.53 (0.18)	0.87 (0.07)	0.92 (0.05)	0.63 (0.17)	0.85 (0.12)	0.94 (0.04)

0.33 to 0.24 to 0.03 for the bond yield. These opposite patterns reflect the opposite signs of the correlations between stock returns and the two predictors, as explained earlier.

The lagged unexpected returns $\{u_s\}_{s\leq t}$ contain a significant amount of information about μ_t beyond that included in x_t. When $\{u_s\}_{s\leq t}$ is added to x_t in estimating μ_t, the R^2s increase by anywhere between 7% and 83%. For example, under the more informative prior on ρ_{uw}, the R^2 increases from 0.03 to 0.86 for the bond yield, from 0.53 to 0.87 for CAY, and from 0.63 to 0.85 when $\{u_s\}_{s\leq t}$ is added to all three predictors. The lagged predictor innovations $\{v_s\}_{s\leq t}$ also contain useful information about μ_t. When $\{v_s\}_{s\leq t}$ is added to x_t and $\{u_s\}_{s\leq t}$, the R^2s increase by between 1% and 41%. The smallest increases, of 1% to 5%, obtain for the dividend yield, while the largest increases, of 9% to 41%, obtain for all three predictors combined.

To summarize, the past values of unexpected returns and predictor innovations contain useful incremental information about the current expected return. This information is used by the predictive system but not by the standard predictive regression.

IV. Conclusion

Predictive systems allow predictors to be imperfectly correlated with the conditional expected return. When predictors are imperfect, expected returns conditional on available data depend not only on the most recent values of those predictors but also on lagged returns and lags of the predictors. Recent returns receive negative weights when a significant portion of the variance in unexpected returns is due to changes in expected returns. The lags of returns and predictors often account for a large fraction of the variation in estimates of conditional expected returns. Predictive systems also allow one to incorporate a prior belief that expected and unexpected returns are negatively correlated. We find that such a belief can have a large impact on estimates of expected returns and various inferences about predictability, including a predictor's correlation with the conditional expected return.

Although our focus is on predictive systems, we also find two implications for predictive regressions. First, we show that if predictors are imperfect, the predictive regression residuals are generally autocorrelated. This autocorrelation should be incorporated when computing standard errors in predictive regressions. In addition, this autocorrelation provides a simple diagnostic for predictor imperfection: Nonzero autocorrelation indicates imperfect predictors. Second, we argue that researchers running predictive regressions should examine the regression-implied correlation between expected and unexpected returns. Predictive regressions in which this correlation is positive are unlikely to perfectly capture time variation in expected stock market returns.

Our initial exploration of predictive systems can be extended in many directions. First, we are intentionally noninformative about the degree of imperfection in a predictor, but one could instead incorporate an informative prior belief about a predictor's correlation with expected return. The latter approach is likely to be preferable when inference is less the objective than is producing the best forecast given one's own prior judgment. Along these lines, one could study the implications of predictive systems for asset allocation. Second, we assume that the conditional mean return follows an AR(1) process, but it would also make sense to consider more complicated processes. For example, if the mean were allowed to have not only a slow-moving persistent component but also a higher-frequency transient component, the bond yield, which is not very persistent, might be inferred to be more highly correlated with the conditional mean. Third, we assume that the return variance is constant, but one could allow it to be time-varying, potentially in a manner correlated with expected return (e.g., Brandt and Kang (2004)). Fourth, we consider three predictors but it would also be interesting to examine the degrees of imperfection in various other predictors that have been proposed in

the literature (e.g., Ferson and Harvey (1991), Lamont (1998), Lewellen (1999), Santos and Veronesi (2006), Ang and Bekaert (2007), etc.). Fifth, we analyze predictability in U.S. stock market returns, but it would also be interesting to apply predictive systems to international markets (e.g., Ferson and Harvey (1993)).

It could also be useful to expand the predictive system to incorporate cash flow news. We have argued that the innovation in the expected return should be negatively correlated with the unexpected return, but if one could account for the portion of the latter that is correlated with cash flow news, the remaining portion would be driven entirely by news about expected return. These issues are beyond the scope of this paper but they merit more attention. See Cochrane (2008), Rytchkov (2007), and van Binsbergen and Koijen (2007) for recent analyses of the interaction between return predictability and cash flow predictability. Cash flow forecasts also enter the calculations of the implied cost of capital, which is used by Pástor, Sinha, and Swaminathan (2008) to proxy for the conditional expected market return.

When compared to predictive regressions, predictive systems typically deliver more precise estimates of expected return. Recall that, in reaching this conclusion, we measure precision as the posterior mean of the ratio of true R^2s. One could instead ask whether expected returns estimated using a predictive system are more precise out of sample, for example, whether they exhibit lower mean squared error (MSE) than a simpler approach such as a predictive regression or the sample average.[21] A Bayesian investor with a quadratic (MSE) loss function would prefer a forecast that combines his priors and the data to estimate the conditional expected return based on the correct model. The correct model, when estimated using a finite sample, tends to produce out-of-sample MSEs higher than those from estimates of simpler models when the true degree of predictability is sufficiently small, as discussed by Clark and West (2006, 2007) and Hjalmarsson (2006). Thus, a simple comparison of out-of-sample MSE's would not speak directly to the question of whether the predictive system is the right model from the investor's perspective. That question, one of model selection, is beyond the scope of this study but could be an interesting area for future research.

Appendix

We summarize here the analytical basis for various relations discussed in the text, and we briefly describe the specification of prior distributions. A lengthier Internet Appendix,[22] available on the *Journal of Finance* web site, provides a more thorough treatment, including details of constructing the prior and posterior distributions.

[21] Goyal and Welch (2003, 2008), Rapach, Strauss, and Zhou (2007), and Campbell and Thompson (2008), etc., investigate the abilities of predictive regressions and sample averages to forecast stock returns out of sample.

[22] The Internet Appendix is available online in the "Supplements and Datasets" section at http://www.afajof.org/supplements.asp

A. General Form of the Predictive System

We define the predictive system in its most general form as a VAR for r_t, x_t, and μ_t, with coefficients restricted so that μ_t is the conditional mean of r_{t+1}. We also assume that x_t and μ_t are stationary with means E_x and E_r. The first-order VAR, for example, is

$$\begin{bmatrix} r_{t+1} - E_r \\ x_{t+1} - E_x \\ \mu_{t+1} - E_r \end{bmatrix} = \begin{bmatrix} 0 & 0 & I \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{bmatrix} \begin{bmatrix} r_t - E_r \\ x_t - E_x \\ \mu_t - E_r \end{bmatrix} + \begin{bmatrix} u_{t+1} \\ v_{t+1} \\ w_{t+1} \end{bmatrix}. \tag{A1}$$

The predictive system in (A1) can be viewed alternatively as simply an unrestricted VAR for returns and predictors when some predictors are unobserved. Specifically, consider an unrestricted VAR for r_t, x_t, and π_t, where π_t has the same dimensions as r_t and contains additional unobserved predictors:

$$\begin{bmatrix} r_{t+1} - E_r \\ x_{t+1} - E_x \\ \pi_{t+1} - E_\pi \end{bmatrix} = \begin{bmatrix} B_{11} & B_{12} & B_{13} \\ B_{21} & B_{22} & B_{23} \\ B_{31} & B_{32} & B_{33} \end{bmatrix} \begin{bmatrix} r_t - E_r \\ x_t - E_x \\ \pi_t - E_\pi \end{bmatrix} + \begin{bmatrix} u_{t+1} \\ v_{t+1} \\ v_{t+1} \end{bmatrix}. \tag{A2}$$

When B_{13} is nonsingular, (A1) and (A2) are equivalent, as shown in the Internet Appendix. The general form in (A1) is treated there. Note also from (A2) that the general form nests the traditional VAR for just r_t and x_t (as the variance of v_{t+1} as well as B_{31}, B_{32}, and B_{33} approach zero). Below we confine attention to the simplified version of the predictive system in equations (3) to (5).

B. Expected Returns and Past Values

Let z_t denote the vector of the observed data at time t, $z_t = (r_t{}'\ x_t{}')'$. Denote the data we observe through time t as $D_t = (z_1, \ldots, z_t)$, and note that our complete data consist of D_T. Let D_0 denote the null information set, so that the unconditional moments are given as

$$\begin{bmatrix} r_t \\ x_t \\ \mu_t \end{bmatrix} \Big| D_0 \sim N \left(\begin{bmatrix} E_r \\ E_x \\ E_r \end{bmatrix}, \begin{bmatrix} V_{rr} & V_{rx} & V_{r\mu} \\ V_{xr} & V_{xx} & V_{x\mu} \\ V_{\mu r} & V_{\mu x} & V_{\mu\mu} \end{bmatrix} \right). \tag{A3}$$

Also define

$$E_z = \begin{bmatrix} E_r \\ E_x \end{bmatrix}, V_{zz} = \begin{bmatrix} V_{rr} & V_{rx} \\ V_{xr} & V_{xx} \end{bmatrix}, V_{z\mu} = \begin{bmatrix} V_{r\mu} \\ V_{x\mu} \end{bmatrix}. \tag{A4}$$

Define the conditional moments

$$a_t = \mathrm{E}(\mu_t \mid D_{t-1}), \qquad\qquad b_t = \mathrm{E}(\mu_t \mid D_t), \qquad\qquad f_t = \mathrm{E}(z_t \mid D_{t-1}), \tag{A5}$$

$$P_t = \text{Var}(\mu_t \mid D_{t-1}), \qquad Q_t = \text{Var}(\mu_t \mid D_t), \qquad R_t = \text{Var}(z_t \mid \mu_t, D_{t-1}),$$

$$\text{(A6)}$$

$$S_t = \text{Var}(z_t \mid D_{t-1}), \qquad G_t = \text{Cov}(z_t, \mu_t \mid D_{t-1}). \tag{A7}$$

Conditioning on the (unknown) parameters of the model is assumed throughout but suppressed in the notation for convenience. As shown in the lengthier Internet Appendix, the above conditional moments can be computed iteratively. First, $a_1 = E_r, P_1 = V_{\mu\mu}, f_1 = E_z, S_1 = V_{zz}, G_1 = V_{z\mu}$,

$$R_1 = S_1 - G_1 P_1^{-1} G_1' \tag{A8}$$

$$Q_1 = P_1 \left(P_1 + G_1' R_1^{-1} G_1 \right)^{-1} P_1 \tag{A9}$$

$$b_1 = a_1 + P_1 \left(P_1 + G_1' R_1^{-1} G_1 \right)^{-1} G_1' R_1^{-1} (z_1 - f_1). \tag{A10}$$

Then, for $t = 2, \ldots, T$,

$$P_t = \beta^2 Q_{t-1} + \Sigma_{ww} \tag{A11}$$

$$S_t = \begin{bmatrix} Q_{t-1} + \Sigma_{uu} & \Sigma_{uv} \\ \Sigma_{vu} & \Sigma_{vv} \end{bmatrix} \tag{A12}$$

$$G_t = \begin{bmatrix} \beta Q_{t-1} + \Sigma_{uw} \\ \Sigma_{vw} \end{bmatrix} \tag{A13}$$

$$R_t = S_t - G_t P_t^{-1} G_t' \tag{A14}$$

$$Q_t = P_t \left(P_t + G_t' R_t^{-1} G_t \right)^{-1} P_t \tag{A15}$$

$$a_t = (1 - \beta) E_r + \beta b_{t-1} \tag{A16}$$

$$f_t = \begin{bmatrix} b_{t-1} \\ (I - A) E_x + A x_{t-1} \end{bmatrix} \tag{A17}$$

$$b_t = a_t + P_t \left(P_t + G_t' R_t^{-1} G_t \right)^{-1} G_t' R_t^{-1} (z_t - f_t) \tag{A18}$$

$$= a_t + G_t' S_t^{-1} (z_t - f_t). \tag{A19}$$

The conditional expected return $b_t = \text{E}(r_{t+1} \mid D_t)$ can be expressed as a function of past returns and predictors. Denote

$$[m_t \; n_t'] \equiv P_t \left(P_t + G_t' R_t^{-1} G_t \right)^{-1} G_t' R_t^{-1} = G_t' S_t^{-1}$$

$$= \text{Cov}(z_t', \mu_t \mid D_{t-1})[\text{Var}(z_t \mid D_{t-1})]^{-1}, \tag{A20}$$

so that, from equation (A18), for $t > 1$,

$$b_t = a_t + [m_t \quad n'_t](z_t - f_t)$$

$$= (1 - \beta)E_r + \beta b_{t-1} + [m_t \quad n'_t] \begin{bmatrix} r_t - b_{t-1} \\ x_t - (I - A)E_x - Ax_{t-1} \end{bmatrix}$$

$$= (1 - \beta)E_r + (\beta - m_t)b_{t-1} + m_t r_t + n'_t v_t, \tag{A21}$$

or

$$b_t - E_r = \beta(b_{t-1} - E_r) + m_t(r_t - b_{t-1}) + n'_t v_t. \tag{A22}$$

For $t = 1$, we obtain

$$b_1 - E_r = m_1(r_1 - b_0) + n'_1 v_1,$$

where v_1 denotes $x_1 - E_x$. Repeated substitution for the lagged values of $(b_t - E_r)$ gives

$$b_t = E_r + \sum_{s=1}^{t} \lambda_s(r_s - b_{s-1}) + \sum_{s=1}^{t} \phi'_s v_s, \tag{A23}$$

where

$$\lambda_s = m_s \beta^{t-s} \tag{A24}$$

$$\phi_s = n_s \beta^{t-s}. \tag{A25}$$

This is equation (8) in the text.

The conditional expected return b_t can be rewritten so that past forecast errors are replaced by returns in excess of the unconditional mean E_r. To do so, modify equation (A21) as

$$b_t - E_r = (\beta - m_t)(b_{t-1} - E_r) + m_t(r_t - E_r) + n'_t v_t, \tag{A26}$$

so that repeated substitution for the lagged values of $(b_t - E_r)$ then yields

$$b_t = E_r + \sum_{s=1}^{t} \omega_s(r_s - E_r) + \sum_{s=1}^{t} \delta'_s v_s, \tag{A27}$$

where

$$\omega_s = \begin{cases} (\beta - m_t)(\beta - m_{t-1}) \cdots (\beta - m_{s+1})m_s, & \text{for } s < t \\ m_s, & \text{for } s = t \end{cases} \tag{A28}$$

$$\delta_s = \begin{cases} (\beta - m_t)(\beta - m_{t-1}) \cdots (\beta - m_{s+1})n_s, & \text{for } s < t \\ n_s, & \text{for } s = t. \end{cases} \tag{A29}$$

This is equation (11) in the text.

If E_r is replaced by the sample mean in equation (11), then the estimate of b_t becomes

$$\hat{b}_t = \sum_{s=1}^{t} \kappa_s r_s + \sum_{s=1}^{t} \delta'_s v_s, \tag{A30}$$

where

$$\kappa_s = \frac{1}{t}\left(1 - \sum_{l=1}^{t} \omega_l\right) + \omega_s, \tag{A31}$$

and $\sum_{s=1}^{t} \kappa_s = 1$. This is equation (26) in the text.

C. R^2 Ratios

The numerator of the R^2 ratio in equation (29) is computed as

$$R^2(\mu_t \text{ on } x_t) = \frac{\text{Var}[\text{E}(\mu_t \mid x_t)]}{\text{Var}(\mu_t)} = \frac{\text{Var}[\text{E}(\mu_t) + V'_{x\mu} V_{xx}^{-1}(x_t - E_x)]}{\text{Var}(\mu_t)} = \frac{V'_{x\mu} V_{xx}^{-1} V_{x\mu}}{\sigma_\mu^2}, \tag{A32}$$

where $\sigma_\mu^2 = \sigma_w^2/(1-\beta^2)$, $V_{x\mu} = (I - \beta A)^{-1}\sigma_{vw}$, and

$$\text{vec}(V_{xx}) = [I - (A \otimes A)]^{-1}\text{vec}(\Sigma_{vv}). \tag{A33}$$

Equation (A33) follows from the relation $V_{xx} = AV_{xx}A' + \Sigma_{vv}$ and the well-known identity $\text{vec}(DFG) = (G' \otimes D)\text{vec}(F)$.

The denominator of the R^2 ratio in equation (29) is computed as

$$R^2(\mu_t \text{ on } D_t) = \frac{\text{Var}[\text{E}(\mu_t \mid D_t)]}{\text{Var}(\mu_t)} = \frac{\text{Var}(\mu_t) - \text{Var}(\mu_t \mid D_t)}{\text{Var}(\mu_t)} = 1 - \frac{Q_t}{\sigma_\mu^2}, \tag{A34}$$

where Q_t is given in equation (A15). We replace Q_t by its steady-state value, Q, which equals a solution of a quadratic equation:

$$Q = \frac{\sqrt{\xi_1^2 - 4\xi_2} - \xi_1}{2}, \tag{A35}$$

$$\xi_1 = (1-\beta^2)\left(\sigma_u^2 - \sigma_{uv}\Sigma_{vv}^{-1}\sigma_{vu}\right) + 2\beta\left(\sigma_{uw} - \sigma_{wv}\Sigma_{vv}^{-1}\sigma_{vu}\right) - \left(\sigma_w^2 - \sigma_{wv}\Sigma_{vv}^{-1}\sigma_{vw}\right)$$

$$= (1-\beta^2)\text{Var}(u \mid v) + 2\beta\text{Cov}(u, w \mid v) - \text{Var}(w \mid v)$$

$$\xi_2 = \left(\sigma_{uw} - \sigma_{wv}\Sigma_{vv}^{-1}\sigma_{vu}\right)^2 - \left(\sigma_u^2 - \sigma_{uv}\Sigma_{vv}^{-1}\sigma_{vu}\right)\left(\sigma_w^2 - \sigma_{wv}\Sigma_{vv}^{-1}\sigma_{vw}\right)$$

$$= \text{Cov}(u, w \mid v)^2 - \text{Var}(u \mid v)\text{Var}(w \mid v) < 0.$$

The value of Q is also used in computing the steady-state values of m_t and n_t from equation (A20),

$$m = (\beta Q + \text{Cov}(u, w \mid v))(Q + \text{Var}(u \mid v))^{-1} \tag{A36}$$

$$n' = (\sigma_{wv} - m\sigma_{uv})\Sigma_{vv}^{-1}. \tag{A37}$$

D. Variance Decomposition of Expected Return

In equation (34), the conditional expected return μ_t depends on three time-varying variables:

1. $C1 = x_t$, the current predictor values,
2. $C2 = \sum_{i=0}^{\infty} \beta^i u_{t-i}$, an infinite sum of current and lagged unexpected returns,
3. $C3 = \sum_{i=0}^{\infty}(\beta^i I_K - A^i)v_{t-i}$, an infinite sum of current and lagged predictor innovations, plus an error term. In the variance decomposition in Table IV, we consider regressions of μ_t on various subsets of $(C1, C2, C3)$. Let C denote a given subset of $(C1, C2, C3)$. The R^2 from the regression of μ_t on C is equal to

$$R^2(\mu_t \text{ on } C) = \frac{V'_{\mu C} V_C^{-1} V_{\mu C}}{\sigma_\mu^2}. \tag{A38}$$

The matrix V_C, the covariance matrix of C, is constructed from

$$\text{Var}(C1) = V_{xx}$$

$$\text{Var}(C2) = \sigma_u^2(1 - \beta^2)^{-1}$$

$$\text{vec}(\text{Var}(C3)) = \left[(1 - \beta^2)^{-1}I_{K^2} - (I_K - \beta A)^{-1} \otimes I_K - I_K \otimes (I_K - \beta A)^{-1}\right.$$
$$\left. + (I_{K^2} - A \otimes A)^{-1}\right]\text{vec}(\Sigma_{vv})$$

$$\text{Cov}(C1, C2) = (I_K - \beta A)^{-1}\sigma_{vu}$$

$$\text{Cov}(C2, C3) = \left[(1 - \beta^2)^{-1}I_K - (I_K - \beta A)^{-1}\right]\sigma_{vu}$$

$$\text{vec}(\text{Cov}(C1, C3')) = \left[I_K \otimes (I_K - \beta A)^{-1} + (I_{K^2} - A \otimes A)^{-1}\right]\text{vec}(\Sigma_{vv}),$$

and $V_{\mu C}$, the vector of covariances between μ_t and C, is constructed from

$$\text{Cov}(\mu_t, C1') = \Psi_v\text{Var}(C1) + \Psi_u\text{Cov}(C1, C2)' + \Psi_v\text{Cov}(C1, C3')'$$

$$\text{Cov}(\mu_t, C2) = \Psi_u\text{Var}(C2) + \Psi_v\text{Cov}(C1, C2) + \Psi_v\text{Cov}(C2, C3)$$

$$\text{Cov}(\mu_t, C3') = \Psi_v\text{Var}(C3) + \Psi_v\text{Cov}(C1, C3') + \Psi_u\text{Cov}(C2, C3)'.$$

E. Prior Distributions

We require both x_t and μ_t to be stationary, so that all eigenvalues of A must lie inside the unit circle and $\beta \in (-1, 1)$. Apart from this restriction, our prior is noninformative about A but informative about β, $\beta \sim N(0.99, 0.15^2)$ (see Figure 5). We put a mildly informative prior on E_r, $E_r \sim N(\bar{\mu}, \sigma_{E_r}^2)$, centered

at the sample mean return with a large prior *SD* of 1% per quarter. We use a noninformative prior for E_x, $E_x \sim N(0, \sigma_{E_x}^2 I_K)$ with a large σ_{E_x}. All four parameters, A, β, E_μ, and E_x, are independent a priori.

The prior on Σ is more complicated. We divide the elements of Σ into two subsets: first, the 2×2 submatrix $\Sigma_{11} \equiv [\sigma_u^2 \, \sigma_{uw}; \sigma_{wu} \, \sigma_w^2]$, and second, the elements of Σ that involve v: $\Sigma_{(v)} \equiv (\Sigma_{vv}, \sigma_{vu}, \sigma_{vw})$. We choose a prior that is informative about Σ_{11} but noninformative about $\Sigma_{(v)}$. Such a prior is obtained as a posterior of Σ when a noninformative prior is updated with a hypothetical sample in which there are T_0 observations of (u, w) but only $S_0 \ll T_0$ observations of v (see Stambaugh (1997)). Details are provided in the lengthier Internet Appendix.

REFERENCES

Ang, Andrew, and Geert Bekaert, 2007, Stock return predictability: Is it there? *Review of Financial Studies* 20, 651–707.

Ang, Andrew, and Monika Piazzesi, 2003, A no-arbitrage vector autoregression of term structure dynamics with macroeconomic and latent variables, *Journal of Monetary Economics* 50, 745–787.

Avramov, Doron, 2002, Stock return predictability and model uncertainty, *Journal of Financial Economics* 64, 423–458.

Avramov, Doron, 2004, Stock return predictability and asset pricing models, *Review of Financial Studies* 17, 699–738.

Avramov, Doron, and Russ Wermers, 2006, Investing in mutual funds when returns are predictable, *Journal of Financial Economics* 81, 339–377.

Baks, Klaas P., Andrew Metrick, and Jessica Wachter, 2001, Should investors avoid all actively managed mutual funds? A study in Bayesian performance evaluation, *Journal of Finance* 56, 45–85.

van Binsbergen, Jules H., and Ralph S. J. Koijen, 2007, Predictive regressions: A present-value approach, Working paper, Duke University.

Brandt, Michael W., and Qiang Kang, 2004, On the relationship between the conditional mean and volatility of stock returns: A latent VAR approach, *Journal of Financial Economics* 72, 217–257.

Campbell, John Y., 1987, Stock returns and the term structure, *Journal of Financial Economics* 18, 373–399.

Campbell, John Y., 1991, A variance decomposition for stock returns, *The Economic Journal* 101, 157–179.

Campbell, John Y., 2001, Why long horizons? A study of power against persistent alternatives, *Journal of Empirical Finance* 8, 459–491.

Campbell, John Y., and John Ammer, 1993, What moves the stock and bond markets? A variance decomposition for long-term asset returns, *Journal of Finance* 48, 3–37.

Campbell, John Y., and Samuel B. Thompson, 2008, Predicting the equity premium out of sample: Can anything beat the historical average? *Review of Financial Studies* 21, 1509–1531.

Campbell, John Y., and Motohiro Yogo, 2006, Efficient tests of stock return predictability, *Journal of Financial Economics* 81, 27–60.

Carter, Chris K., and Robert Kohn, 1994, On Gibbs sampling for state space models, *Biometrika* 81, 541–553.

Casella, G., and E. I. George, 1992, Explaining the Gibbs sampler, *The American Statistician* 46, 167–174.

Cavanagh, Christopher L., Graham Elliott, and James H. Stock, 1995, Inference in models with nearly integrated regressors, *Econometric Theory* 11, 1131–1147.

Clark, Todd E., and Kenneth D. West, 2006, Using out-of-sample mean squared prediction errors to test the martingale difference hypothesis, *Journal of Econometrics* 135, 155–186.

Clark, Todd E., and Kenneth D. West, 2007, Approximately normal tests for equal predictive accuracy in nested models, *Journal of Econometrics* 138, 291–311.

Cochrane, John H., 2008, The dog that did not bark: A defense of return predictability, *Review of Financial Studies* 21, 1533–1575.

Conrad, Jennifer, and Gautam Kaul, 1988, Time-variation in expected returns, *Journal of Business* 61, 409–425.

Cremers, K. J. Martijn, 2002, Stock return predictability: A Bayesian model selection perspective, *Review of Financial Studies* 15, 1223–1249.

Dangl, Thomas, and Michael Halling, 2006, Equity return prediction: Are coefficients time-varying?, Working paper, University of Vienna.

Duffee, Gregory R., 2006, Are variations in term premia related to the macroeconomy? Working paper, University of California, Berkeley.

Elliott, Graham, and James H. Stock, 1994, Inference in time series regression when the order of integration of a regressor is unknown, *Econometric Theory* 10, 672–700.

Fama, Eugene F., and Kenneth R. French, 1988, Dividend yields and expected stock returns, *Journal of Financial Economics* 22, 3–26.

Fama, Eugene F., and G. William Schwert, 1977, Asset returns and inflation, *Journal of Financial Economics* 5, 115–146.

Ferson, Wayne E., and Campbell R. Harvey, 1991, The variation of economic risk premiums, *Journal of Political Economy* 99, 385–415.

Ferson, Wayne E., and Campbell R. Harvey, 1993, The risk and predictability of international equity returns, *Review of Financial Studies* 6, 527–566.

Ferson, Wayne E., Sergei Sarkissian, and Timothy T. Simin, 2003, Spurious regressions in financial economics? *Journal of Finance* 58, 1393–1413.

Frühwirth-Schnatter, Sylvia, 1994, Data augmentation and dynamic linear models, *Journal of Time Series Analysis* 15, 183–202.

Goyal, Amit, and Ivo Welch, 2003, Predicting the equity premium with dividend ratios, *Management Science* 49, 639–654.

Goyal, Amit, and Ivo Welch, 2008, A comprehensive look at the empirical performance of equity premium prediction, *Review of Financial Studies* 21, 1455–1508.

Hamilton, James D., 1994, *Time Series Analysis* (Princeton University Press, Princeton, NJ).

Harvey, Andrew C., 1989, *Forecasting, Structural Time Series Models and the Kalman Filter* (Cambridge University Press, Cambridge, UK).

Hjalmarsson, Erik, 2006, Should we expect significant out-of-sample results when predicting stock returns? Working paper, Board of Governors of the Federal Reserve System.

Jansson, Michael, and Marcelo J. Moreira, 2006, Optimal inference in regressions with nearly integrated regressors, *Econometrica* 74, 681–714.

Johannes, Michael, Nicholas Polson, and Jon Stroud, 2002, Sequential optimal portfolio performance: Market and volatility timing, Working paper, Columbia University.

Jones, Christopher S., and Jay Shanken, 2005, Mutual fund performance with learning across funds, *Journal of Financial Economics* 78, 507–552.

Kandel, Shmuel, and Robert F. Stambaugh, 1987, Long-horizon returns and short-horizon models, Working paper, University of Chicago.

Kandel, Shmuel, and Robert F. Stambaugh, 1996, On the predictability of stock returns: An asset allocation perspective, *Journal of Finance* 51, 385–424.

Keim, Donald B., and Robert F. Stambaugh, 1986, Predicting returns in the stock and bond markets, *Journal of Financial Economics* 17, 357–390.

Kothari, S. P., Jonathan Lewellen, and Jerold B. Warner, 2006, Stock returns, aggregate earnings surprises, and behavioral finance, *Journal of Financial Economics* 79, 537–568.

Lamont, Owen, 1998, Earnings and expected returns, *Journal of Finance* 53, 1563–1587.

Lamoureux, Christopher G., and Guofu Zhou, 1996, Temporary components of stock returns: What do the data tell us? *Review of Financial Studies* 9, 1033–1059.

Lettau, Martin, and Sydney C. Ludvigson, 2001, Consumption, aggregate wealth, and expected stock returns, *Journal of Finance* 56, 815–849.

Lettau, Martin, and Sydney C. Ludvigson, 2005, Expected returns and dividend growth, *Journal of Financial Economics* 76, 583–626.

Lettau, Martin, and Stijn Van Nieuwerburgh, 2008, Reconciling the return predictability evidence, *Review of Financial Studies* 21, 1607–1652.

Lewellen, Jonathan, 1999, The time-series relations among expected return, risk, and book-to-market, *Journal of Financial Economics* 54, 5–43.

Lewellen, Jonathan, 2004, Predicting returns with financial ratios, *Journal of Financial Economics* 74, 209–235.

Mankiw, N. Gregory, and Matthew D. Shapiro, 1986, Do we reject too often?: Small sample properties of tests of rational expectations models, *Economics Letters* 20, 139–145.

Menzly, Lior, Tano Santos, and Pietro Veronesi, 2004, Understanding predictability, *Journal of Political Economy* 112, 1–47.

Nelson, Charles R., Myung J. Kim, 1993, Predictable stock returns: The role of small sample bias, *Journal of Finance* 48, 641–661.

Pástor, Ľuboš, 2000, Portfolio selection and asset pricing models, *Journal of Finance* 55, 179–223.

Pástor, Ľuboš, Meenakshi Sinha, and Bhaskaran Swaminathan, 2008, Estimating the intertemporal risk-return tradeoff using the implied cost of capital, *Journal of Finance* 63, 2859–2897.

Pástor, Ľuboš, and Robert F. Stambaugh, 1999, Costs of equity capital and model mispricing, *Journal of Finance* 54, 67–121.

Pástor, Ľuboš, and Robert F. Stambaugh, 2000, Comparing asset pricing models: An investment perspective, *Journal of Financial Economics* 56, 335–381.

Pástor, Ľuboš, and Robert F. Stambaugh, 2001, The equity premium and structural breaks, *Journal of Finance* 56, 1207–1239.

Pástor, Ľuboš, and Robert F. Stambaugh, 2002a, Mutual fund performance and seemingly unrelated assets, *Journal of Financial Economics* 63, 315–349.

Pástor, Ľuboš, and Robert F. Stambaugh, 2002b, Investing in equity mutual funds, *Journal of Financial Economics* 63, 351–380.

Rapach, David E., Jack K. Strauss, and Guofu Zhou, 2007, Out-of-sample equity premium prediction: Consistently beating the historical average, *Review of Financial Studies* (forthcoming).

Rozeff, Michael S., 1984, Dividend yields are equity risk premiums, *Journal of Portfolio Management* 11(1), 68–75.

Rytchkov, Oleg, 2007, Filtering out expected dividends and expected returns, Working paper, MIT.

Santos, Tano, and Pietro Veronesi, 2006, Labor income and predictable stock returns, *Review of Financial Studies* 19, 1–44.

Shiller, Robert J., 1981, Do stock prices move too much to be justified by subsequent changes in dividends? *American Economic Review* 71, 421–436.

Stambaugh, Robert F., 1986, Bias in regressions with lagged stochastic regressors, Working paper, University of Chicago.

Stambaugh, Robert F., 1997, Analyzing investments whose histories differ in length, *Journal of Financial Economics* 45, 285–331.

Stambaugh, Robert F., 1999, Predictive regressions, *Journal of Financial Economics* 54, 375–421.

Wachter, Jessica, and Missaka Warusawitharana, 2006, Predictable returns and asset allocation: Should a skeptical investor time the market? Working paper, University of Pennsylvania.

Zellner, Arnold, 1971, *An Introduction to Bayesian Inference in Econometrics* (John Wiley & Sons, New York).